

Web caches/ cookies and the DNS

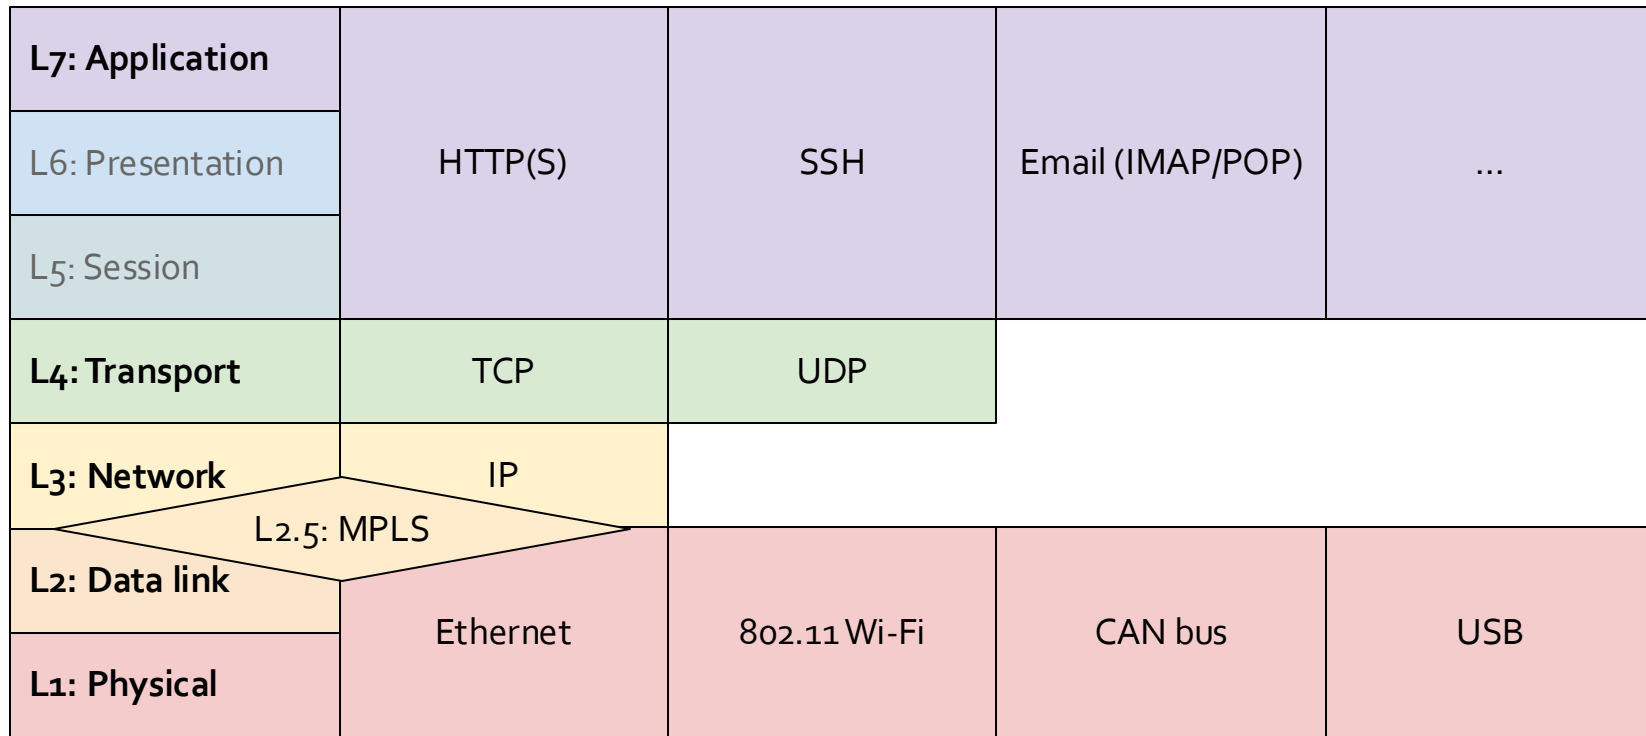
CE 352, Computer Networks

Salem Al-Agtash

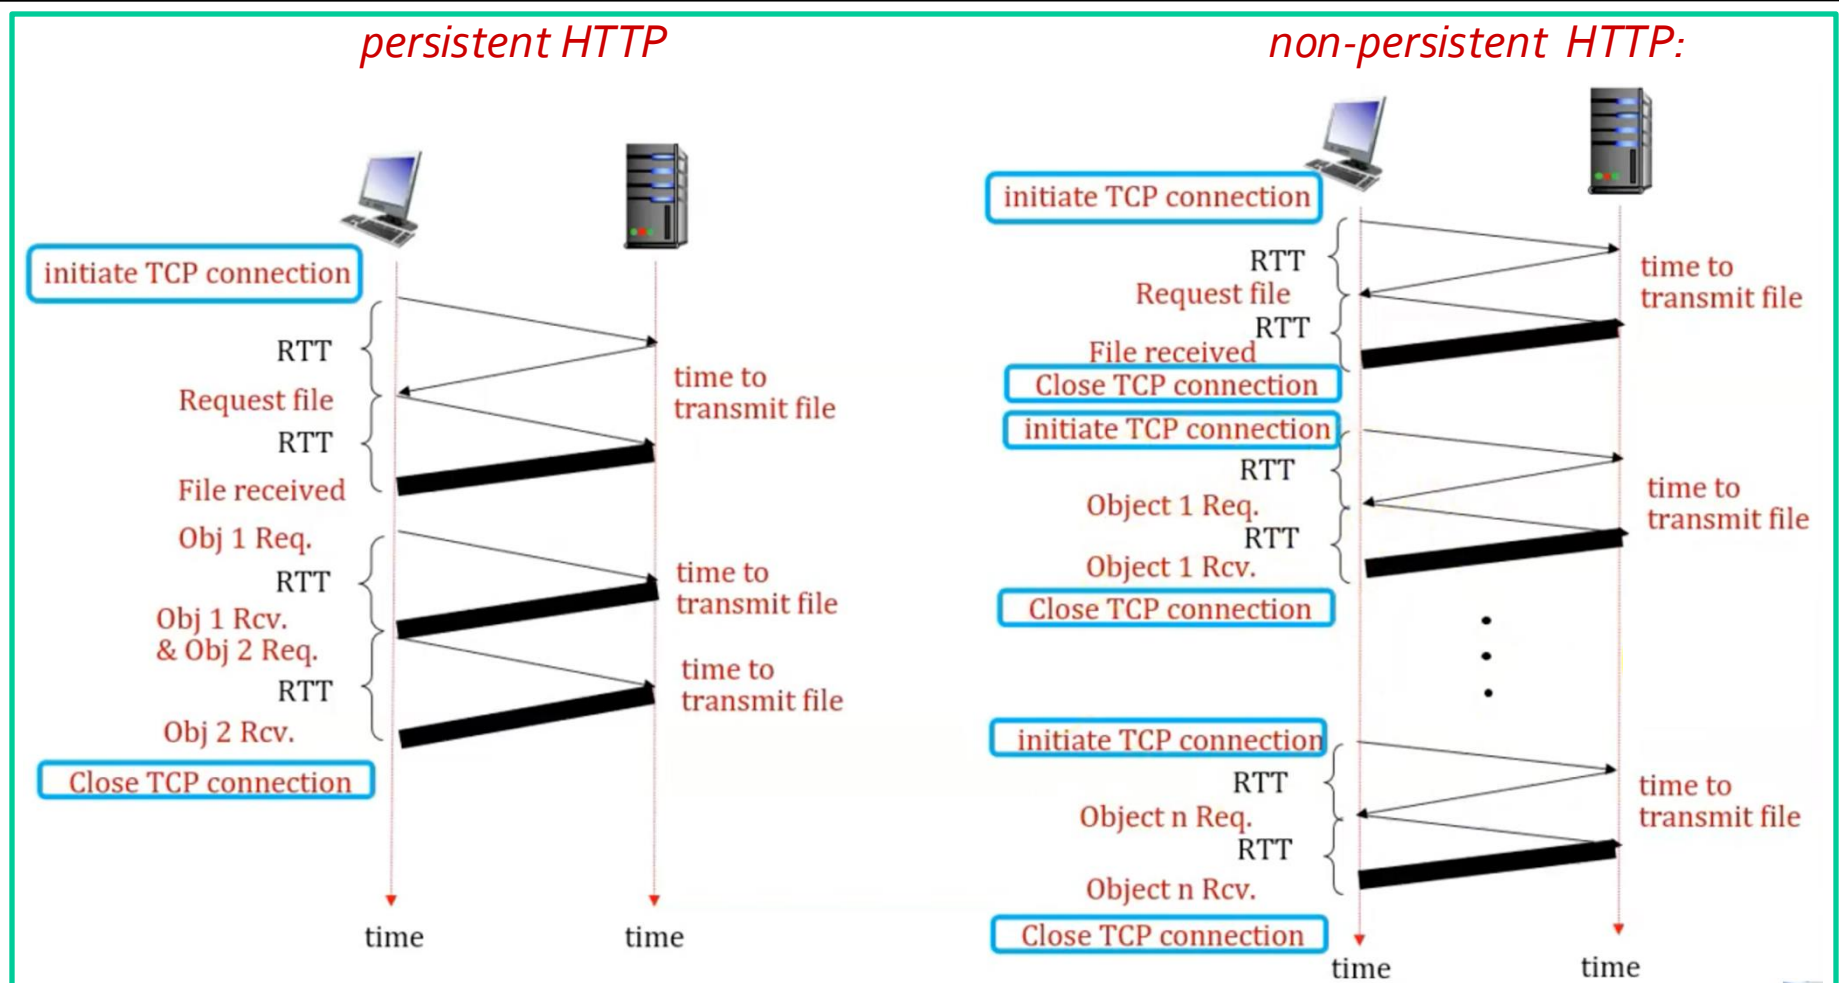
Lecture 5

Slides are adapted from Computer Networking: A Top Down Approach, 7th Edition © J.F Kurose and K.W. Ross

Recap (Internet Protocol Stack)



Recap (Persistent vs no-persistent http)



Topics of today

- Part1

- 1- Stateless and stateful web servers
- 2- Cookies to maintain state in stateless protocol
- 3- Uses of cookies
- 4- Web caches and caching examples

- Part2

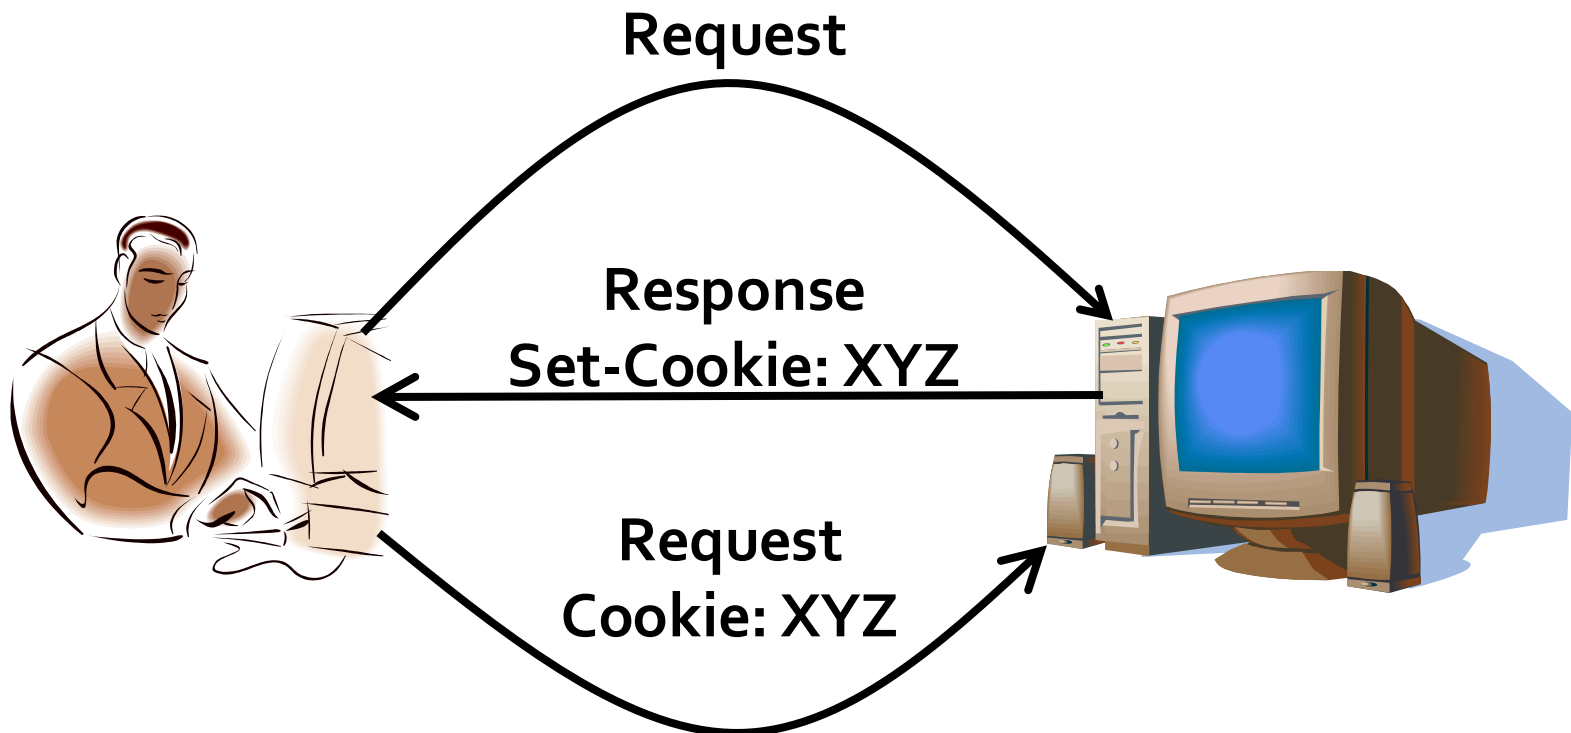
- 1- What is a DNS and its hierarchy
- 2- The DNS 13 root name servers
- 3- The TLD authoritative servers
- 4- The local DNS name server
- 5- The DNS name resolution example
- 6- DNS records: host and dig commands
- 7- DNS caching and updating records
- 8- DNS properties and attacks

Cookies: State in a Stateless Protocol

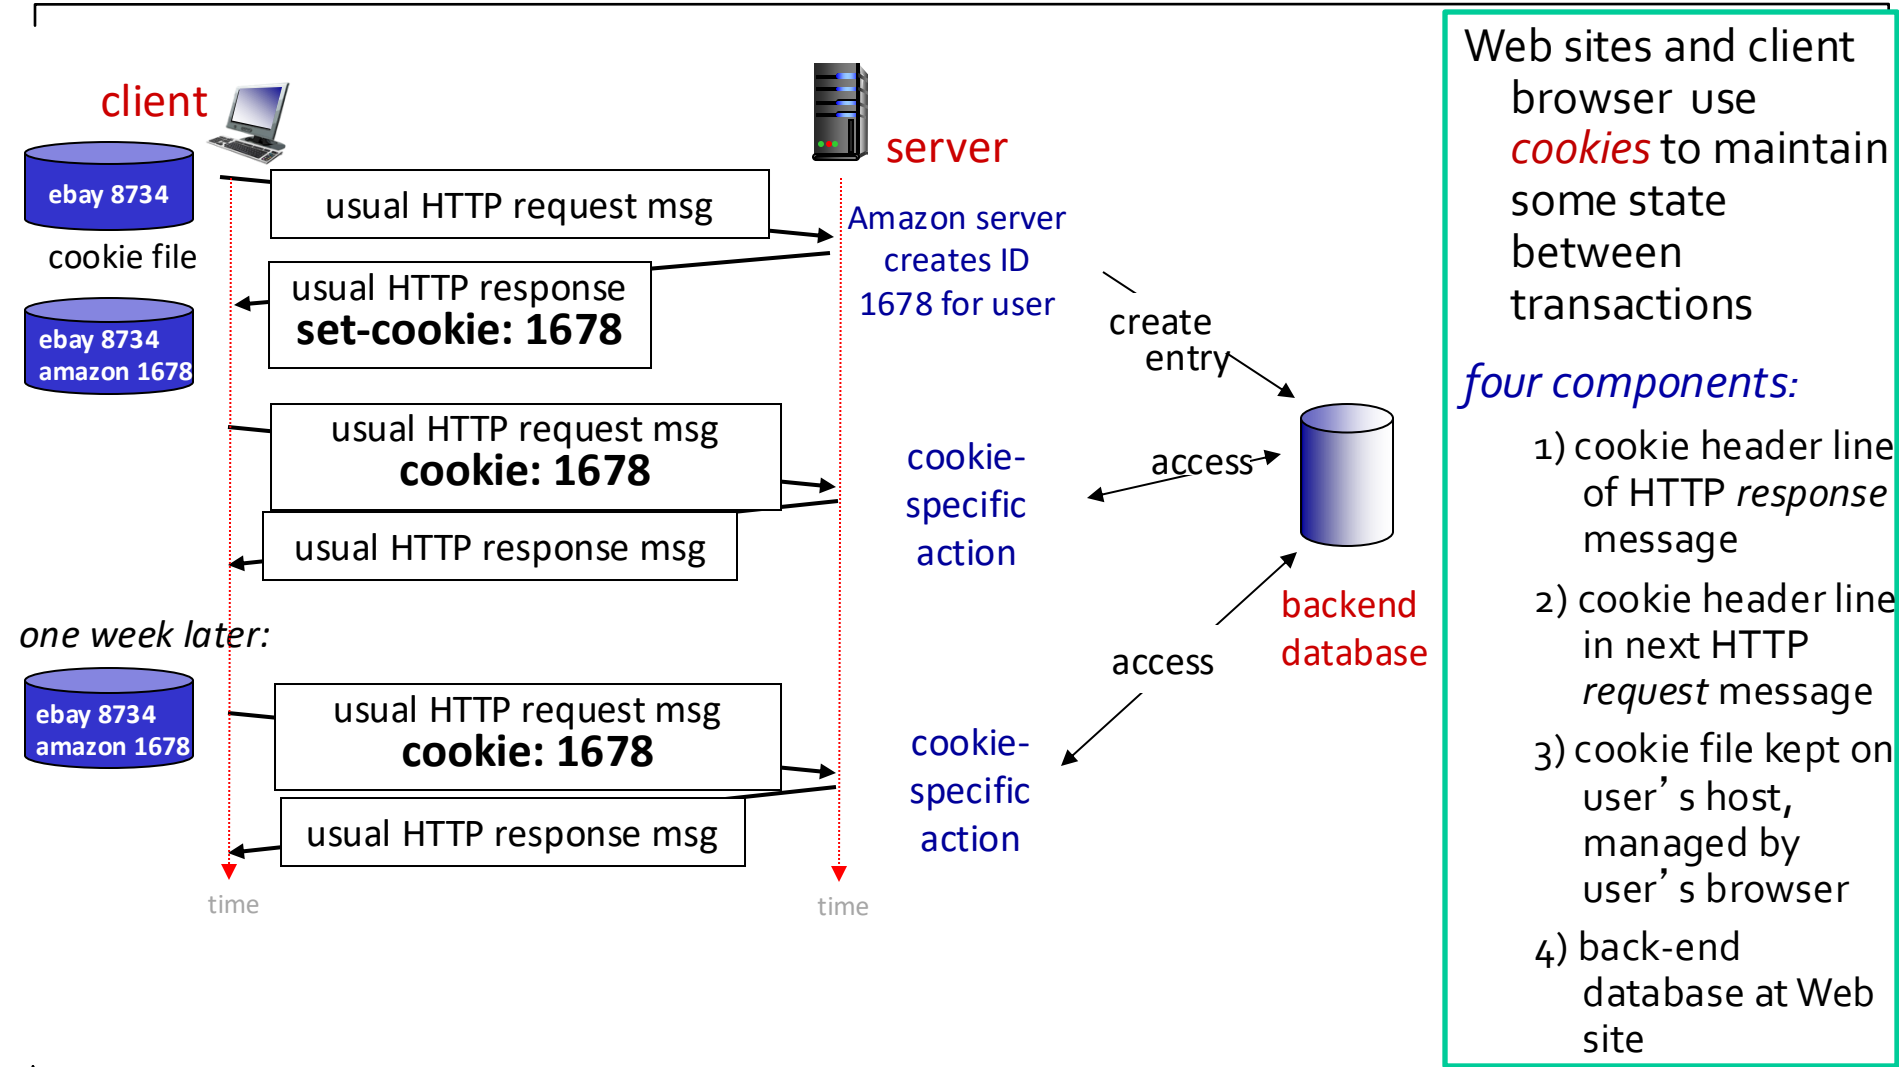
Client-side state maintenance

- ▣ Client stores small state on behalf of server
- ▣ Client sends state in future requests to the server

Can provide authentication



Cookies: keeping “state”



Use of cookies

what cookies can be used for:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

aside

cookies and privacy:


- cookies permit sites to learn a lot about you
- third party persistent cookies (tracking cookies) allow common identity (cookie value) to be tracked across multiple web sites to sites

Web caches (proxy server)

- Responding to client requests without involving origin server
- Web users
 - Availability and fast downloads
- Web content providers
 - More users, cost-effective infrastructure, and non-congested network

Why Web caching?

- reduce response time for client request
 - cache is closer to client
- reduce traffic on an institution's access link
- Internet is dense with caches
 - enables "poor" content providers to more effectively deliver content



- Proxies: Caching and replication
- Content delivery networks (CDN): economies of scale


Caching example:

assumptions:

- avg request rate from browsers to origin servers: 1.5 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

consequences:

- LAN utilization:?
- access link utilization:?
- total delay =



Caching example:

assumptions:

- avg request rate from browsers to origin servers: 1.5 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps


consequences:

- LAN utilization: $(1.5 \text{ Mbps}) / 1 \text{ Gbps} = 0.15\%$
- access link utilization: $1.5 \text{ Mbps} / 1.54 \text{ Mbps} = 97\%$
- total delay = Internet delay + access delay + LAN delay

Assuming local access 0.01 sec, then: total delay = $2 + 0.1 = 2.01$ sec


As access link approaches 100% utilization, link upgrade is needed for users to get better internet response!

This delay becomes very large when the link approaches 100% utilization, unacceptable for institution's users. So: Either upgrade the link or install a local cache



Install local cache


- suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
 - 60% of requests use access link
- data rate to browsers over access link = ?
- utilization = ?
- total delay = ?



Install local cache

- suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
 - 60% of requests use access link
- data rate to browsers over access link
 - $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
- utilization = $0.9 / 1.54 = 58\%$
- total delay
 - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$

Assuming local access 0.01 sec, then: total delay = $0.6 (2.01) + 0.4 (0.1) = 1.21 \text{ sec}$
So with no link upgrade, users still get better internet response!



The DNS: domain name system

Internet hosts, routers:

- IP address (32 or 128 bits) - used for addressing datagrams
- “name”, e.g., www.yahoo.com - used by humans

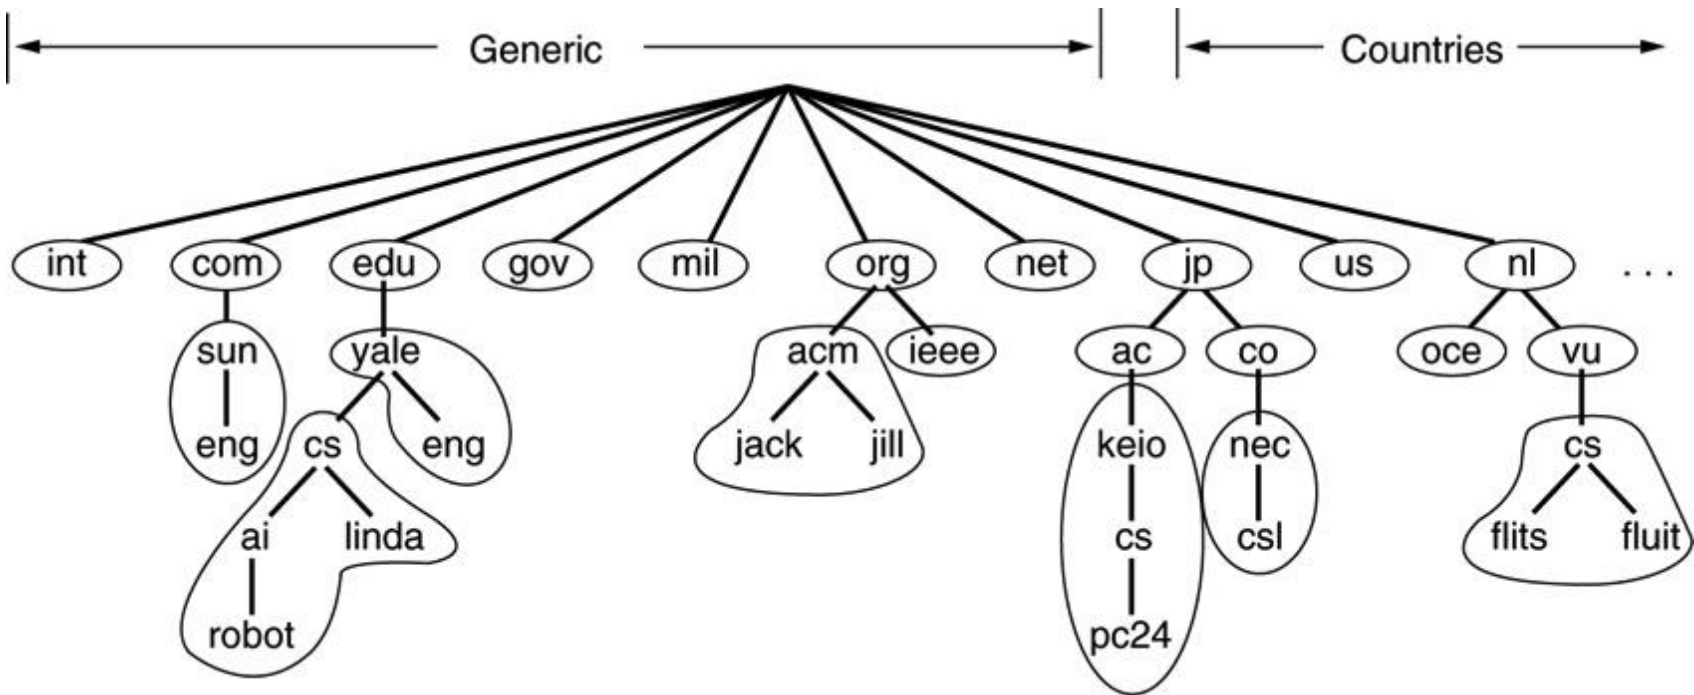
Domain Name System:

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol*: hosts, name servers communicate to *resolve* names (address/name translation)

handles many trillions of queries/day:


- *many* more reads than writes
- *performance matters*: almost every Internet transaction interacts with DNS - msec count!

DNS hierarchy



- DNS name space is divided into zones.
- Each zone contains some part of the tree and contains name servers holding information about the that zone

DNS: a distributed, hierarchical database



client wants IP for www.amazon.com; 1st approximation:

- client queries amazon.com DNS server to get IP address for www.amazon.com
- client queries .com DNS server to get amazon.com DNS server
- client queries root server to find .com DNS server

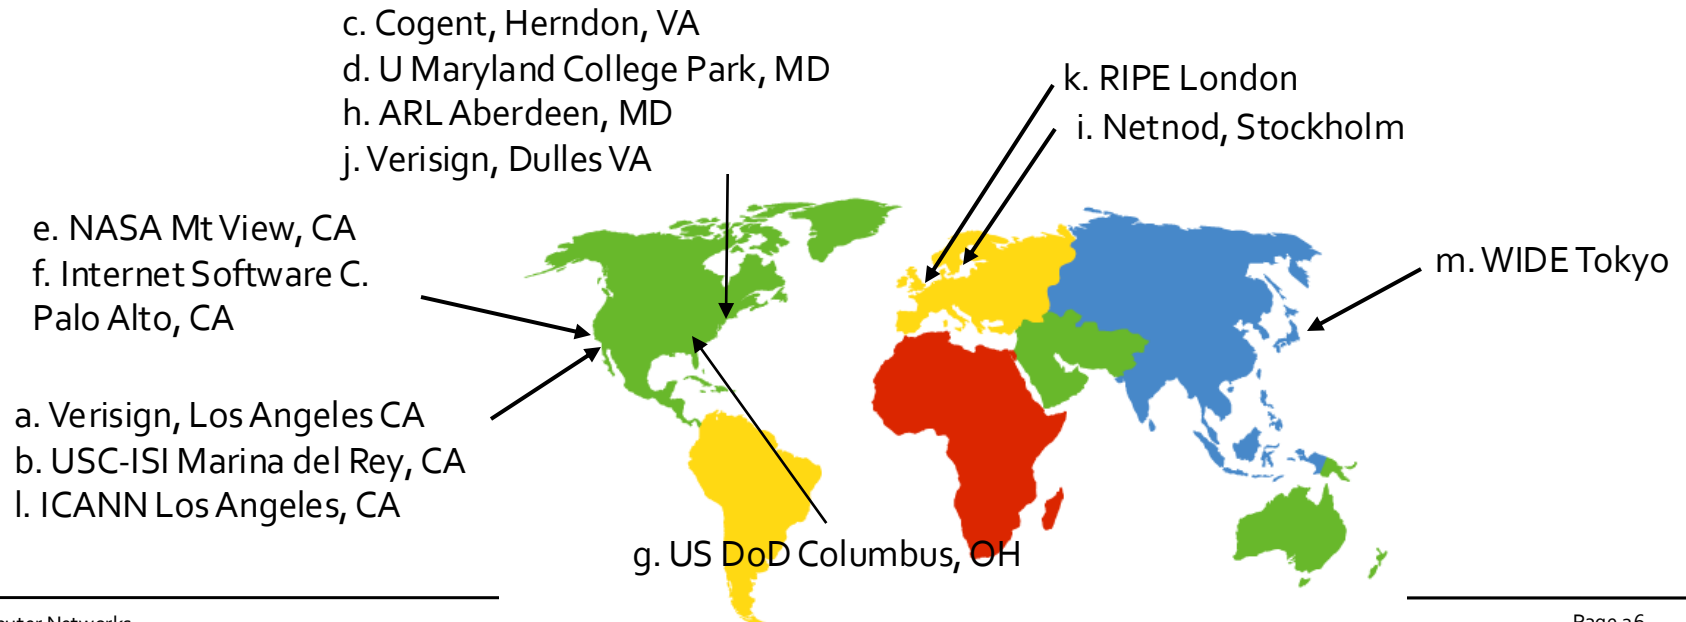
DNS: root name servers (1)

13 root servers (labeled A – M) <http://www.root-servers.org/>

contacted by local name server that can not resolve name

root name server:


- contacts authoritative name server if name mapping not known
- gets mapping
- returns mapping to local name server



TLD (2), authoritative (3) servers

top-level domain (TLD) servers:

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD




authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

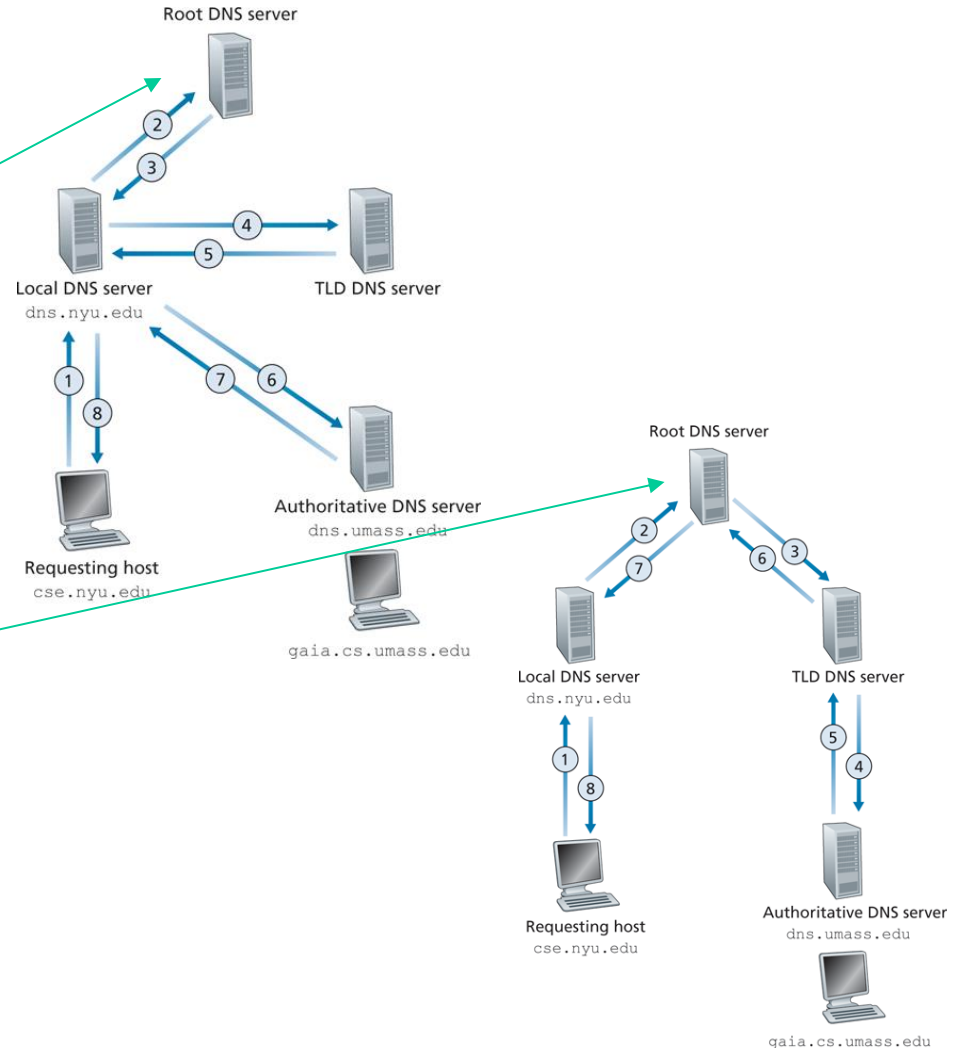
Local DNS name (4) server

- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
- has local cache of recent name-to-address translation pairs (but may be out of date!)
- acts as proxy, forwards query into hierarchy



DNS name resolution example

- Client-server interaction on UDP Port 53
- host at cse.nyu.edu wants IP address for gaia.cs.umass.edu
- *iterative query*:
 - contacted server replies with name of server to contact
 - “I don’t know this name, but ask this server”
 - Widely used in practice
- *recursive query*:
 - puts burden of name resolution on contacted name server
 - heavy load at upper levels of hierarchy?



DNS records

DNS: distributed database storing resource records (RR)

RR format: (name, value, type, ttl)

Type	Meaning	Value
SOA	Start of Authority	Parameters for this zone
A	IP address of a host	32-Bit integer
MX	Mail exchange	Priority, domain willing to accept e-mail
NS	Name Server	Name of a server for this domain
CNAME	Canonical name	Domain name
PTR	Pointer	Alias for an IP address
HINFO	Host description	CPU and OS in ASCII
TXT	Text	Uninterpreted ASCII text

host – a www.gju.edu.jo

```
$host -a www.gju.edu.jo
Trying "www.gju.edu.jo"
;; ->HEADER<- opcode: QUERY, status: NOERROR, id: 32587
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 1
```

```
;; QUESTION SECTION:
;www.gju.edu.jo. IN ANY
```

```
;; ANSWER SECTION:
www.gju.edu.jo. 1140 IN A 87.236.233.242
```

```
;; AUTHORITY SECTION:
gju.edu.jo. 1543 IN NS dns1.junet.edu.jo.
gju.edu.jo. 1543 IN NS mail-relay.gju.edu.jo.
gju.edu.jo. 1543 IN NS ns1.gju.edu.jo.
```

```
;; ADDITIONAL SECTION:
ns1.gju.edu.jo. 1753 IN A 87.236.233.230
```

```
Received 132 bytes from fe80::1%5#53 in 7 ms
```

Try:

```
dig www.gju.edu.jo
dig www.gju.edu.jo any
dig www.gju.edu.jo MX
dig www.gju.edu.jo SOA
```

Inserting records into DNS

- Example: new startup “Network Utopia”
- Get a block of space from ISP
 - e.g. 220.44.8.128/25
- Register networkuptopia.com at *DNS registrar* (e.g., Network Solutions)
 - provide registrar with names and IP addresses of authoritative name server (primary and secondary)
 - registrar inserts two RRs into .com TLD server:
 - (networkuptopia.com, dns1.networkuptopia.com, NS)
 - (dns1.networkuptopia.com, 220.44.8.129, A)
- create authoritative server dns1.networkuptopia.com:
 - type A record for www.networkuptopia.com;
 - type MX record for networkuptopia.com

DNS: caching, updating records

- Performing DNS queries takes time
- Caching can reduce overhead
 - once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time "*time to live*" (TTL)
- TLD servers typically cached in local name servers thus root name servers not often visited
 - cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- update/notify mechanisms proposed IETF standard RFC 2136

DNS properties and attacks

- Important properties of DNS

- Easy unique naming and reasonable trust model
- Caching lends scalability and performance

- “Distributed Denial of Service” DDoS attacks

- bombard root servers with traffic
 - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers, potentially more dangerous

- redirect attacks

- man-in-middle, Intercept queries
- DNS poisoning
 - Send fake records to DNS server, which caches

- exploit DNS for DDoS

- send queries with spoofed source address: target IP
- requires amplification

Summary

Today:

- Web caches and cookies
- Caching example
- The DNS
- DNS records
- DNS properties and attacks

Canvas discussion:

- Reflection
- Exit ticket

Next time:

- Read 2.3 and 2.5 of K&R
- follow on Canvas! Material and announcements

Any questions?