

# The eMail, P2P applications, CDN

CE 352, Computer Networks

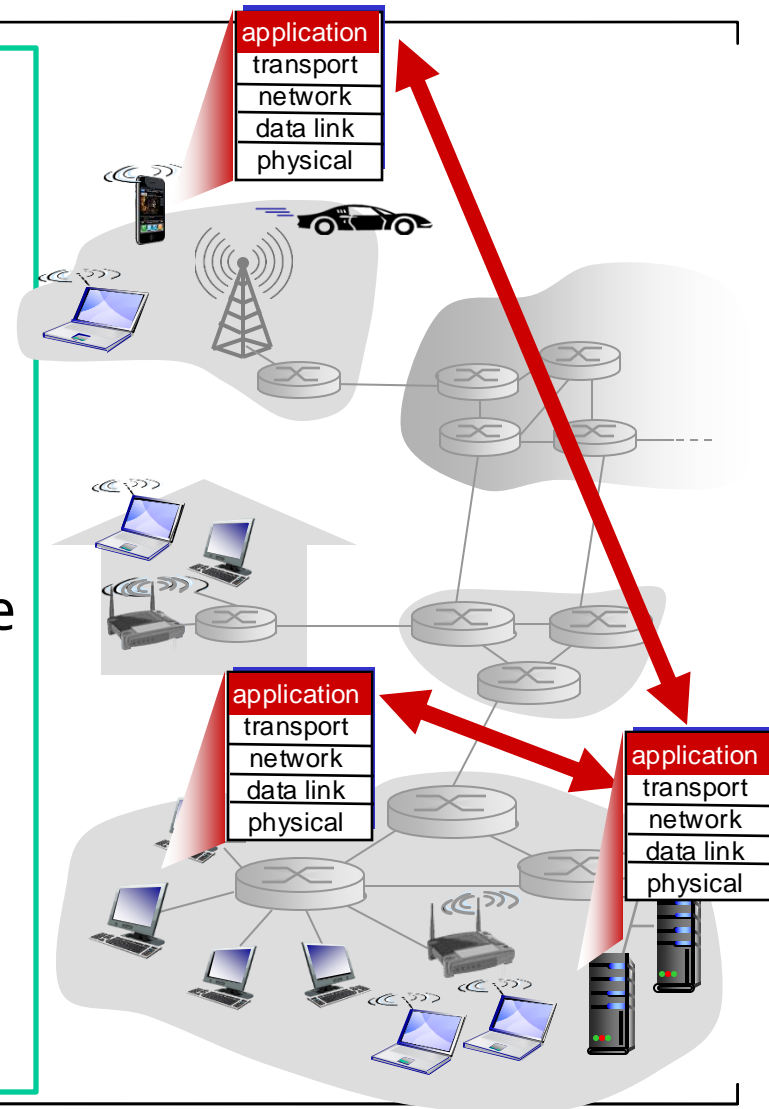
Salem Al-Agtash

Lecture 6

Slides are adapted from Computer Networking: A Top Down Approach, 7<sup>th</sup> Edition © J.F Kurose and K.W. Ross

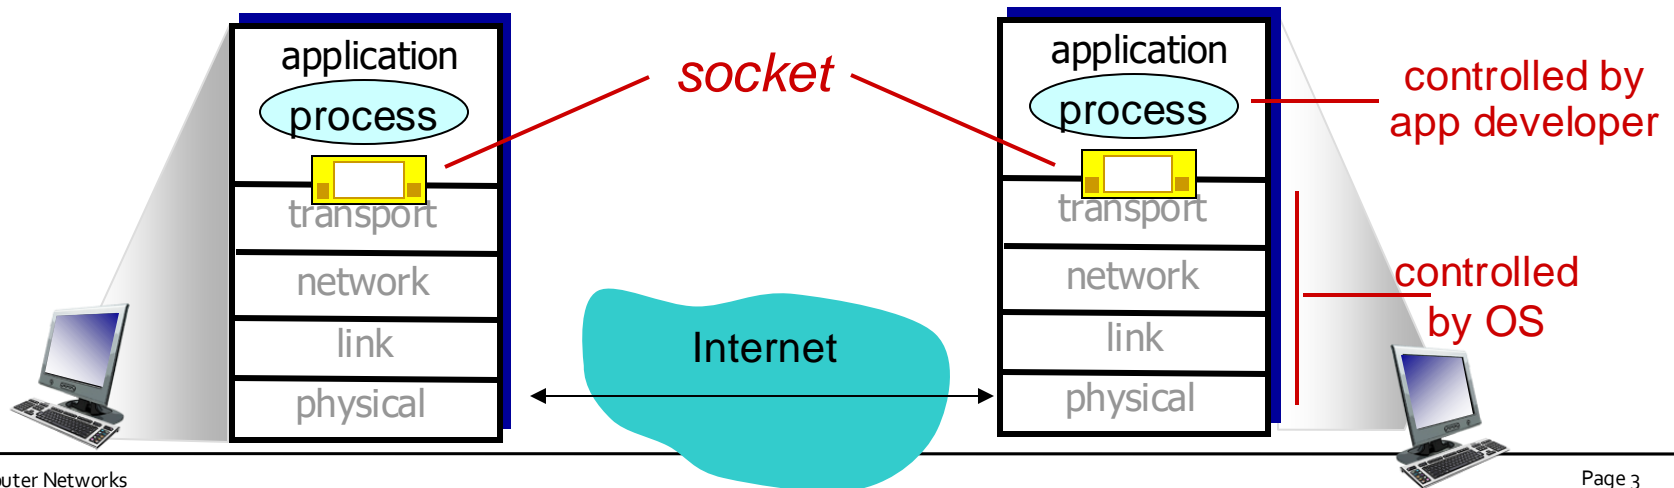
# Recap (Applications on the Network)

- End-end system programs
- Architecture
  - Client – Server
  - Peer-to-peer (P2P)
- no need to write software for network-core devices
- Examples:
  - Web, e-mail, text messaging, remote login, file transfer
  - social networking, multi-user network games
  - VoIP, streaming stored video (YouTube, Hulu, Netflix)




# Recap (process communication)

- Process: Program in Execution
  - Same hosts: processes communicate using IPC defined by OS. e.g. Pipes, Shared Memory, Message Queues
  - Different hosts: processes communicate by exchanging messages. e.g. Client-Server, P2P
- Socket: Process sends/ receives messages via **socket (IP + Port)**
  - Sending process shoves message out door and relies on transport infrastructure to deliver message to socket at receiving process




# Datagram socket: UDP)

- UDP: no “connection” between client & server
  - no handshaking before sending data
  - sender explicitly attaches IP destination address and port # to each packet
  - receiver extracts sender IP address and port# from received packet
- UDP: transmitted data may be lost or received out-of-order
- Application viewpoint: UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server



# Stream socket: TCP

- TCP: “connection” between client & server
  - Server creates socket and begins to listen
  - Client contacts server by creating TCP socket, specifying IP/port of server
  - Server creates thread to communication with particular client
- **Application viewpoint:** TCP provides *reliable* in-order byte stream transfer (“pipe”) between client and server



# System calls

## Fill in IP and Port

- `struct sockaddr_in servAddr, clientAddr;`

## Create a socket

- `socket(AF_INET, SOCK_STREAM, 0);`

## Bind the socket

- `bind(sockfd, (...) &servAddr, sizeof(servAddr))`

## Server listens for connections

- `listen(sockfd, n);`

## Client connects to a server


- `connect(sockfd, (...sockaddr*) &servAddr, sizeof(servAddr));`

## Sever accepts connection

- `accept(sockfd, (struct sockaddr *) &clientAddr, sizeof(clientAddr));`

## Read/write, send/receive

## Binding address to socket




# Recap (WWW)

- Distributed database of “pages” linked through HTTP
  - HTTP/0.9, 1.0, 1.1, 2.0
  - Persistent and non-persistent HTTP, request, response
  - DNS (UDP) and HTTP (TCP)
  - HTML and Dynamic web pages
- Web components
  - Infrastructure:
    - Clients, Servers
  - Content:
    - URL: naming content
    - HTML: formatting content
  - Protocol for exchanging information: HTTP, HTTPS
- Cookies and web caches (proxy server)

# Recap (Web caches (proxy server))

- Web users
  - Availability and fast downloads
- Web content providers
  - More users, cost-effective infrastructure, and non-congested network



- Proxies: Caching and replication
- Content delivery networks (CDN): economies of scale



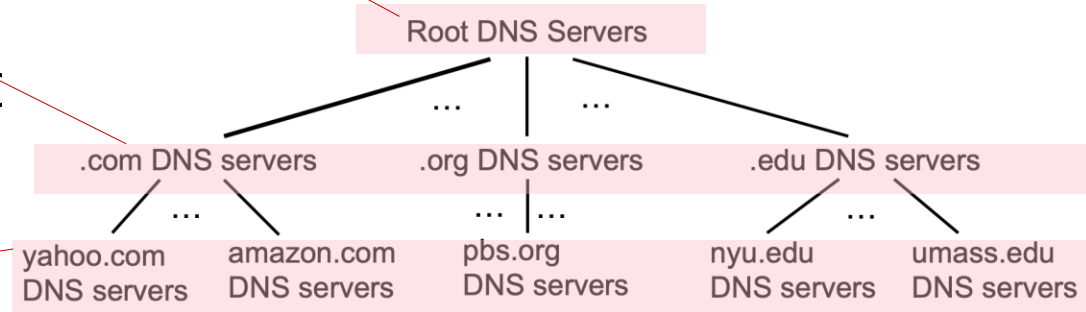
# Recap (DNS: a distributed, hierarchical database)

## *13 root servers (labeled A – M):*

- contacts authoritative name server if name mapping not known
- returns mapping to local name server

## *top-level domain (TLD) servers:*

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca
- Network Solutions for .com TLD
- Educause for .edu TLD




## *authoritative DNS servers:*


- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

# Recap (DNS)

- Client--server on UDP Port 53
- DNS hierarchy
  - Root DNS server
  - TLD DNS server
  - Local/ authoritative DNS server
- DNS (BIND, Windows DNS)
  - Hostname → IP
  - Host and mail server aliasing
  - Load distribution
- DNS records
  - Caching (root server not often visited)
- DNS attacks
  - DDoS and DNS poisoning



*Iterative query*



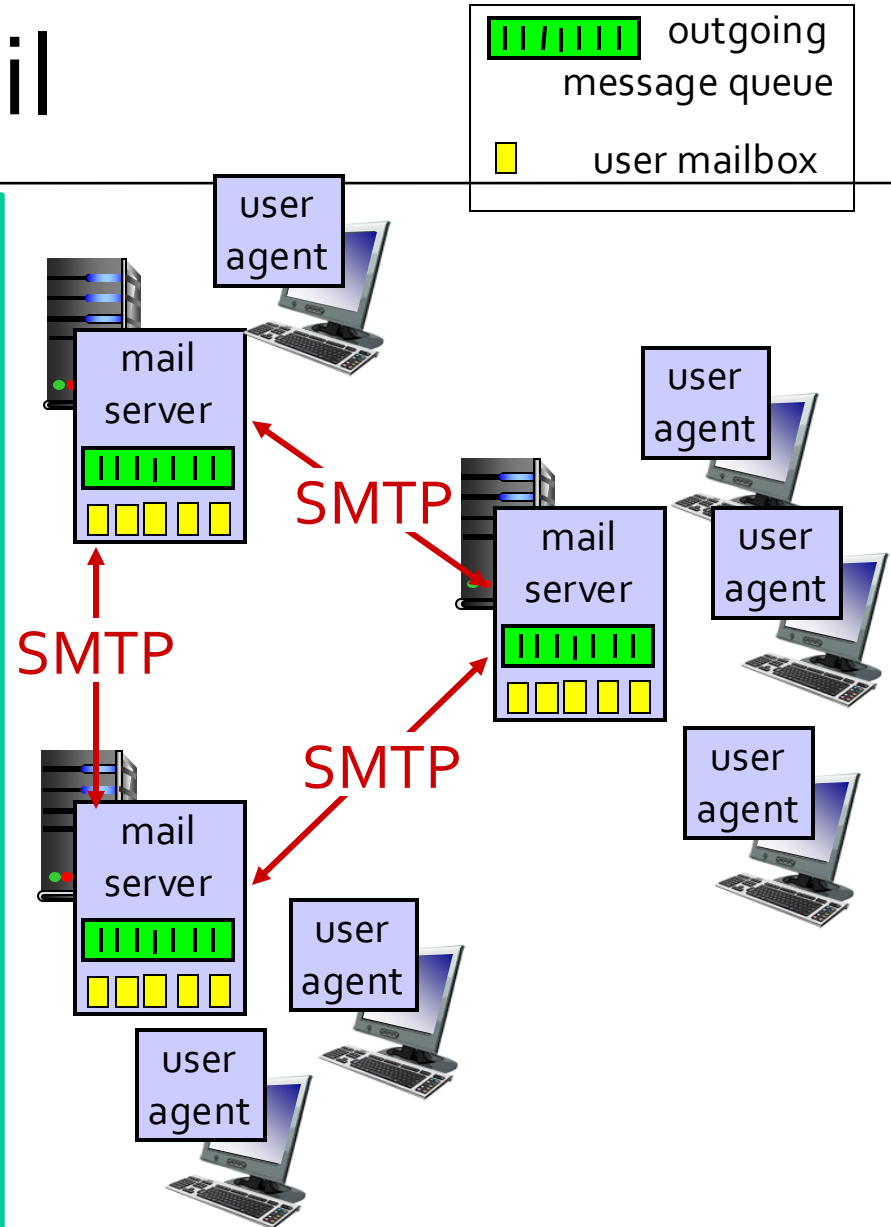
*Recursive query*

- Today

- electronic mail: SMTP, POP<sub>3</sub>, IMAP
- P2P applications
- video streaming and content distribution networks (CDNs)

# Component of eMail

- **User Agent** a.k.a. “mail reader”
  - composing, editing, reading mail messages
  - e.g., Outlook, Thunderbird, Mail
  - outgoing, incoming messages stored on server
- **Mail servers:**
  - *mailbox* contains incoming messages for user
  - *message queue* of outgoing (to be sent) mail messages
- **SMTP protocol**
  - between mail servers
  - client: sending mail server
  - “server”: receiving mail server



# SMTP [RFC 2821]

uses TCP to reliably transfer email message from client to server, port 25

direct transfer: sending server to receiving server

three phases of transfer

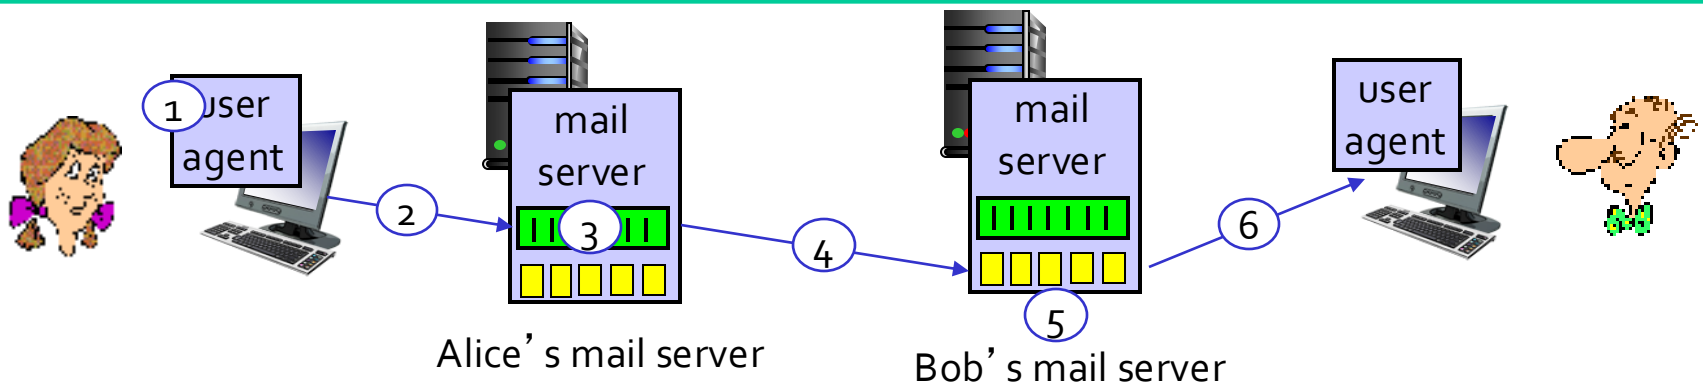
- ▣ handshaking (greeting)
- ▣ transfer of messages
- ▣ closure

command/response interaction (like HTTP)

- ▣ **commands:** ASCII text
- ▣ **response:** status code and phrase

# Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message “to” bob@some school.edu
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



# Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

- *TRY Yourself*

```
telnet servername
25
```

see 220 reply from server  
enter HELO, MAIL FROM,  
RCPT TO, DATA, QUIT  
commands

above lets you send email  
without using email client  
(reader)

# Sample SMTP interaction

```
Connection closed by foreign host.  
ENGR-L-00688:~ salagtach$ telnet mail.smtp2go.com 2525  
Trying 173.255.233.87...  
Connected to mail.smtp2go.com.  
Escape character is '^]'.  
220 mail.smtp2go.com ESMTP Exim 4.94.2-S2G Wed, 19 Jan 2022 00:38:43 +0000  
EHLO  
250-mail.smtp2go.com Hello [162.229.186.16]  
250-SIZE 52428800  
250-8BITMIME  
250-DSN  
250-PIPELINING  
250-PIPE_CONNECT  
250-AUTH CRAM-MD5 PLAIN LOGIN  
250-CHUNKING  
250-STARTTLS  
250-PRDR  
250-SMTPUTF8  
250 HELP  
HELO  
250 mail.smtp2go.com Hello [162.229.186.16]  
MAIL FROM:<alagtach@gmail.com>  
250 OK
```

- *TRY Yourself*

**telnet servername  
25**


see 220 reply from server  
enter HELO, MAIL FROM,  
RCPT TO, DATA, QUIT  
commands

above lets you send email  
without using email client  
(reader)

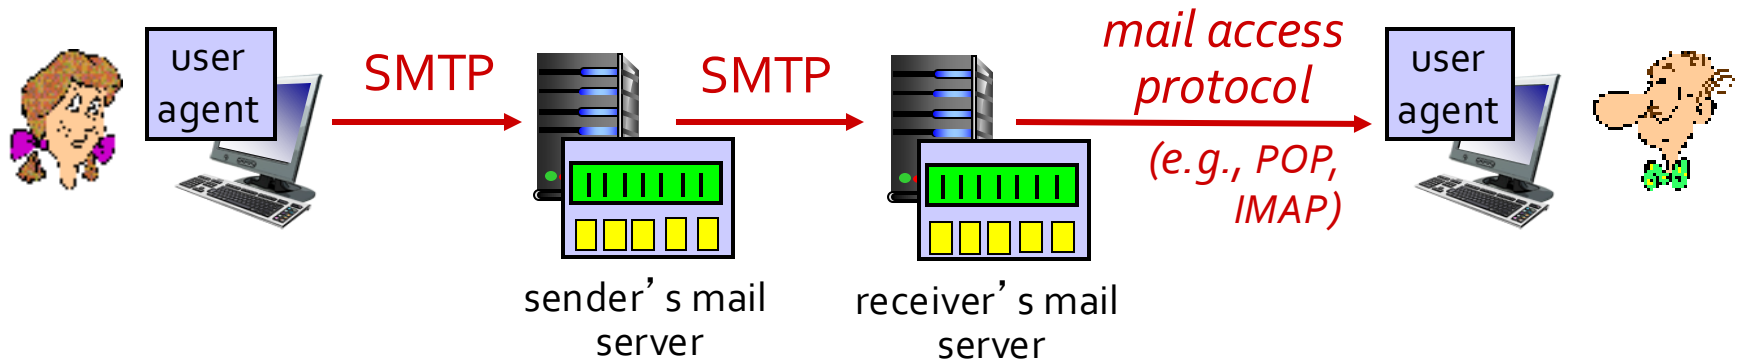


# Mail message format

- SMTP: protocol for exchanging email messages
- RFC 822: standard for text message format:
  - header lines, e.g.,
    - From: alice@crepes.fr
    - To: bob@hamburger.edu
    - Subject: Searching for the meaning of life.
- Body: the “message”
- ASCII characters only



# Mail access protocols



- **SMTP**: delivery/storage to receiver's server
- Mail access protocol: retrieval from server
  - **POP**: Post Office Protocol [RFC 1939]: authorization, download – port 110
  - **IMAP**: Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored messages on server - 143
  - **HTTP**: gmail, Hotmail, Yahoo! Mail, etc.

## Mail settings Network Solutions

POP / IMAP	pop3
Incoming server	mail.[domain].com
Incoming port	995
SSL (security) incoming	ssl
Outgoing server	smtp.[domain].com
Outgoing port	587
Requires sign-in	yes

# POP3 protocol

## *authorization phase*

client commands:

- ▣ **user:** declare username
- ▣ **pass:** password

server responses

- ▣ **+OK**
- ▣ **-ERR**

## *transaction phase*

client:

- ▣ **list:** list message numbers
- ▣ **retr:** retrieve message by number
- ▣ **dele:** delete
- ▣ **quit**

S: +OK POP3 server ready  
C: user bob  
S: +OK  
C: pass hungry  
S: +OK user successfully logged on

C: list  
S: 1 498  
S: 2 912  
S: .  
C: retr 1  
S: <message 1 contents>  
S: .  
C: dele 1  
C: retr 2  
S: <message 1 contents>  
S: .  
C: dele 2  
C: quit

S: +OK POP3 server signing off

# POP<sub>3</sub> and IMAP

## *more about POP<sub>3</sub>*

previous example uses POP<sub>3</sub>  
“download and delete”  
mode

- Bob cannot re-read e-mail if he changes client

POP<sub>3</sub> “download-and-keep”:  
copies of messages on  
different clients

POP<sub>3</sub> is stateless across  
sessions

## *IMAP*

keeps all messages in one  
place: at server

allows user to organize  
messages in folders

keeps user state across  
sessions:


- names of folders and mappings between message IDs and folder name

# Pure P2P architecture

- ❑ *not* always-on server
- ❑ arbitrary end systems directly communicate
- ❑ peers are intermittently connected and change IP addresses

## *examples:*

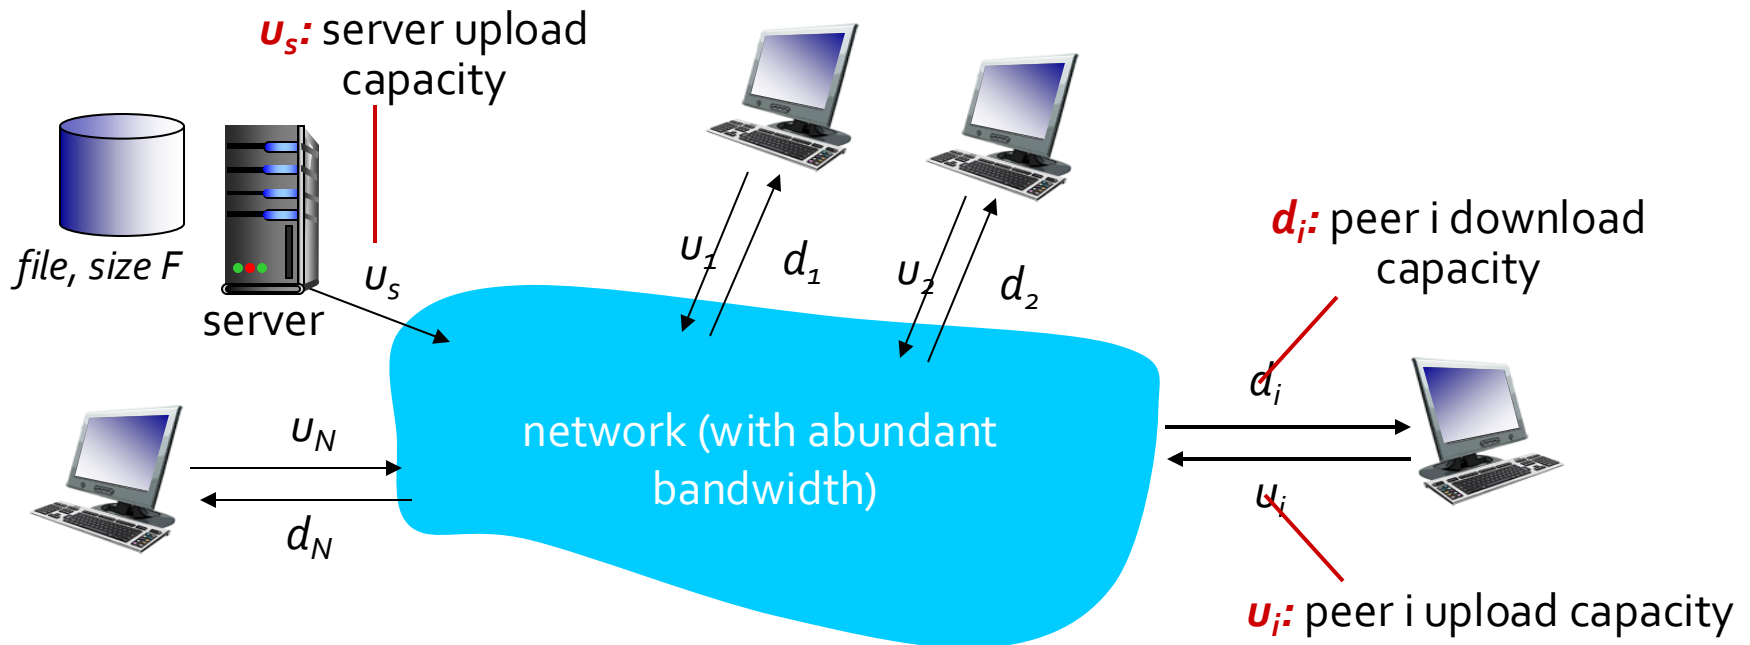
- ❑ file distribution (BitTorrent)
- ❑ Streaming (KanKan)
- ❑ VoIP (Skype)



# File distribution: client-server vs P2P

Time to distribute file (size  $F$ ) from one server to  $N$  peers?


- peer upload/download
- capacity is limited resource



# File distribution time: client-server

*server transmission*: must sequentially send (upload)  $N$  file copies:


- ▣ time to send one copy: ?
- ▣ time to send  $N$  copies: ?
- ▣ *client*: each client must download file copy
- ▣  $d_{min}$  = min client download rate
- ▣ client download time (min  $d$ ): ?



*time to distribute  $F$   
to  $N$  clients using  
client-server approach* ?

# File distribution time: client-server

- server transmission*: must sequentially send (upload)  $N$  file copies:
- time to send one copy:  $F/u_s$
  - time to send  $N$  copies:  $NF/u_s$
  - *client*: each client must download file copy
  - $d_{min}$  = min client download rate
  - client download time (min  $d$ ):  $F/d_{min}$



*time to distribute  $F$   
to  $N$  clients using  
client-server approach*

$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in  $N$



# File distribution time: P2P

*server transmission:* must upload at least one copy


- ▣ time to send one copy: ?

*client:* each client must download file copy

- ▣ client download time with min  $d$ : ?

*clients:* as aggregate must upload  $NF$  bits

- ▣ max upload rate: ?
- ▣ min distrib time: ?



*time to distribute  $F$   
to  $N$  clients using  
P2P approach* ?

# File distribution time: P2P

*server transmission:* must upload at least one copy


- time to send one copy:  $F/u_s$

*client:* each client must download file copy

- client download time with mid  $d$ :  $F/d_{\min}$

*clients:* as aggregate must upload  $NF$  bits

- max upload rate:  $u_s + \text{Sum}(u_i)$
- min distrib time:  $NF/[u_s + \text{Sum}(u_i)]$



*time to distribute  $F$   
to  $N$  clients using  
P2P approach*

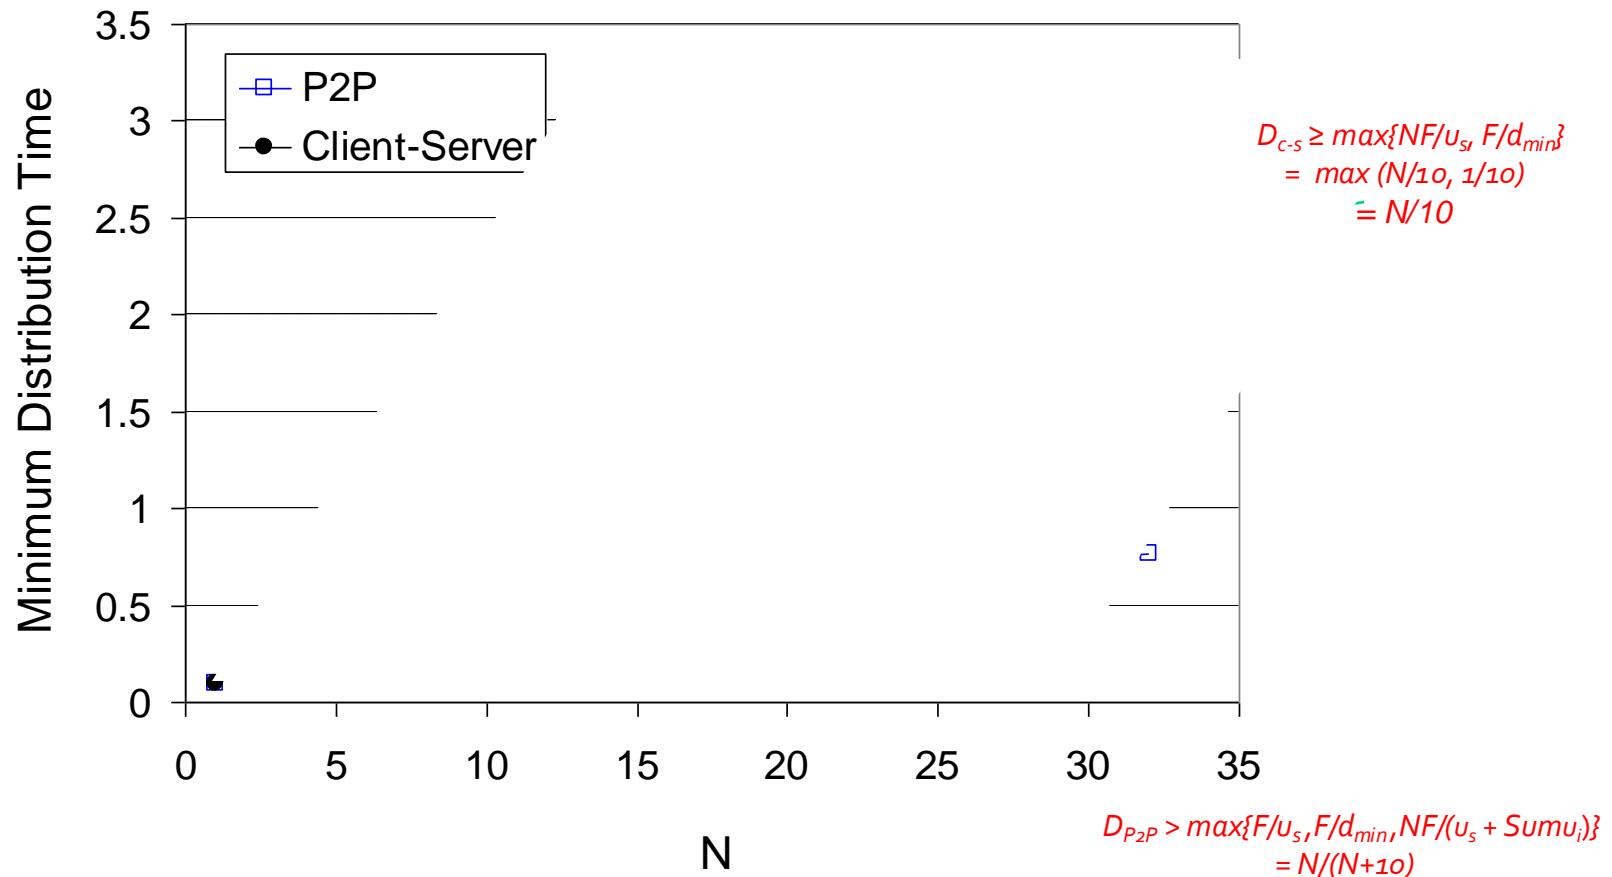
$$D_{P2P} > \max\{F/u_s, F/d_{\min}, NF/(u_s + \text{Sum}u_i)\}$$

increases linearly in  $N$  ...

... but so does this, as each peer brings service capacity


# Client-server vs. P2P: example

- client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$



# Client-server vs. P2P: example

- client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$




# P2P file distribution: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks

*tracker*: tracks peers participating in torrent

*torrent*: group of peers exchanging chunks of a file




Alice arrives ...  
... obtains list  
of peers from tracker

... and begins exchanging  
file chunks with peers in torrent using TCP

# P2P file distribution: BitTorrent

peer joining torrent:

- ❑ has no chunks, but will accumulate them over time from other peers
- ❑ registers with tracker to get list of peers, connects to subset of peers (“neighbors”)

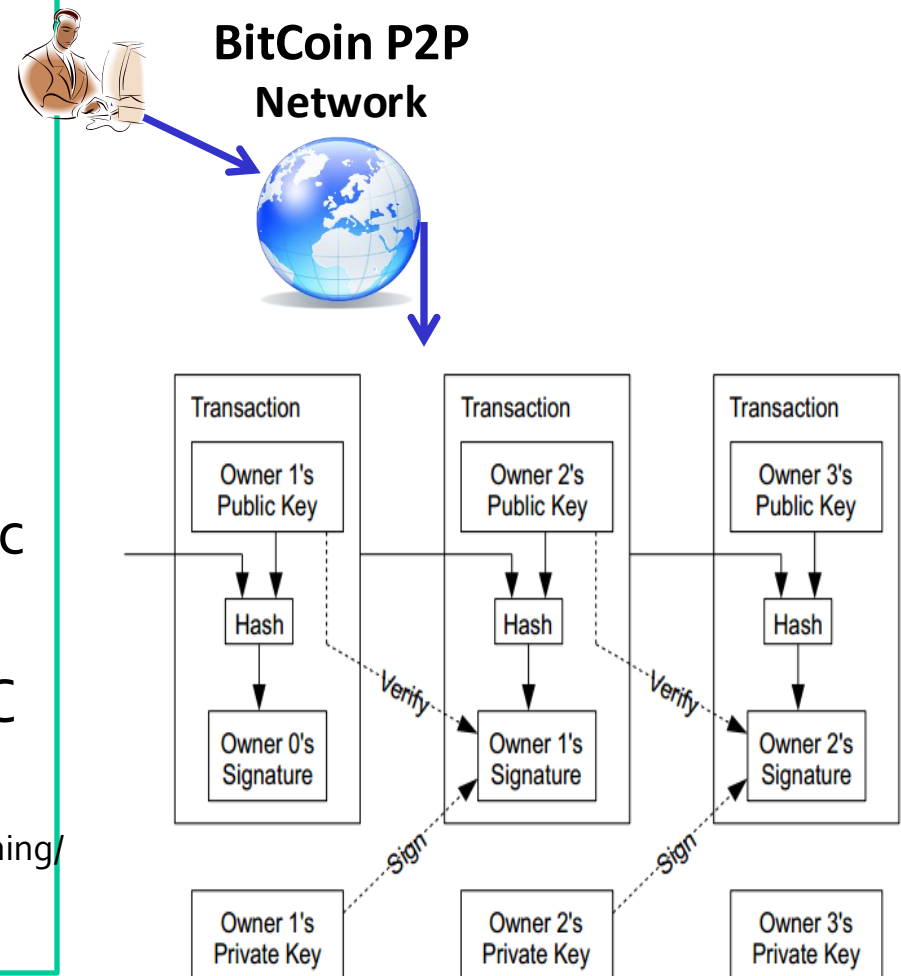


- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- *churn*: peers may come and go
- once peer has entire file, it may (selfishly) leave or (voluntarily) remain in torrent

# P2P Bitcoin mining

- Peer-to-peer computer process used to secure and verify bitcoin transactions
- A distributed ledger that tracks fund transfers among accounts
- Bitcoin transfer:  $\text{Sign}(\text{Previous transaction} + \text{New owner's public key})$
- Mining one bitcoin with just a PC now takes millions of years!

<http://www.techtangerine.com/2017/09/22/bitcoin-pc-mining/>



# Video Streaming and CDNs: context

- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
  - single mega-video server won't work (why?)
- challenge: heterogeneity
  - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution*: distributed, application-level infrastructure






# Multimedia: video

- video: sequence of images displayed at constant rate
  - e.g., 24 images/sec
- digital image: array of pixels
  - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits to encode image
  - spatial (within image)
  - temporal (from one image to next)
- CBR: (constant bit rate) for video encoding
- VBR: (variable bit rate) for video encoding
- examples:
  - MPEG 1 (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)
  - Low-quality (100Kbps), HD (3Mbps), 4K (10 Mbps)

*spatial coding example:*

instead of sending  $N$  values of same color (all purple), send only two values: color value (purple) and number of repeated values ( $N$ )



frame  $i$


*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



frame  $i+1$

# Streaming stored video:

- HTTP streaming: all clients receive same encoding of video
- But, clients have different amount of bandwidth available
- Solution: videos are encoded in different bit-rate versions, different quality videos. So clients dynamically requests chunks of video segments based on bandwidth availability -> DASH



***DASH: Dynamic, Adaptive Streaming over HTTP***

# Streaming multimedia: DASH

## *server:*

- ▣ divides video file into multiple chunks
- ▣ each chunk stored, encoded at different rates
- ▣ *manifest file*: provides URLs for different chunks

## *client:*

- ▣ periodically measures server-to-client bandwidth
- ▣ consulting manifest, requests one chunk at a time
  - ▣ chooses maximum coding rate sustainable given current bandwidth
  - ▣ can choose different coding rates at different time (depending on bandwidth)

## *"intelligence" at client:* client determines

- ▣ *when* to request chunk (so that buffer starvation, or overflow does not occur)
- ▣ *what encoding rate* to request (higher quality when more bandwidth available)
- ▣ *where* to request chunk (can request from URL server that is "close" to client or has high available bandwidth)

# Content distribution networks

*challenge:* how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

*option 1:* single, large “mega-server”

- ❑ single point of failure
- ❑ point of network congestion
- ❑ long path to distant clients
- ❑ multiple copies of video sent over outgoing link

....quite simply: this solution *doesn't scale*

# CDN: Content distribution networks

*challenge:* how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?

*option 2:* store/serve multiple copies of videos at multiple geographically distributed sites (*CDN*). Two placement scenarios:


- ▣ *enter deep:* push CDN servers deep into many access networks
  - ▣ close to users
  - ▣ used by Akamai, 1700 locations
- ▣ *bring home:* smaller number (10's) of larger clusters in POPs near (but not within) access networks -> similar to web caching
  - ▣ used by Limelight

# Google's network infrastructure

Three tiers of server clusters (2011, 2016 data):


- 14 “mega data centers,” each having ~ 100,000 servers, responsible for serving dynamic (and often personalized) content, including search results and Gmail messages.
  - 8 in North America,
  - 4 in Europe
  - 2 in Asia
- ~50 clusters in IXPs scattered throughout the world, each having 100–500 servers, responsible for serving static content, including YouTube videos.
- 100s of “enter-deep” clusters located within an access ISP, each having 10s of servers, responsible to perform TCP splitting and serve static content

# Netflix stores copies of MadMen on CDN nodes. Subscriber → near CDN



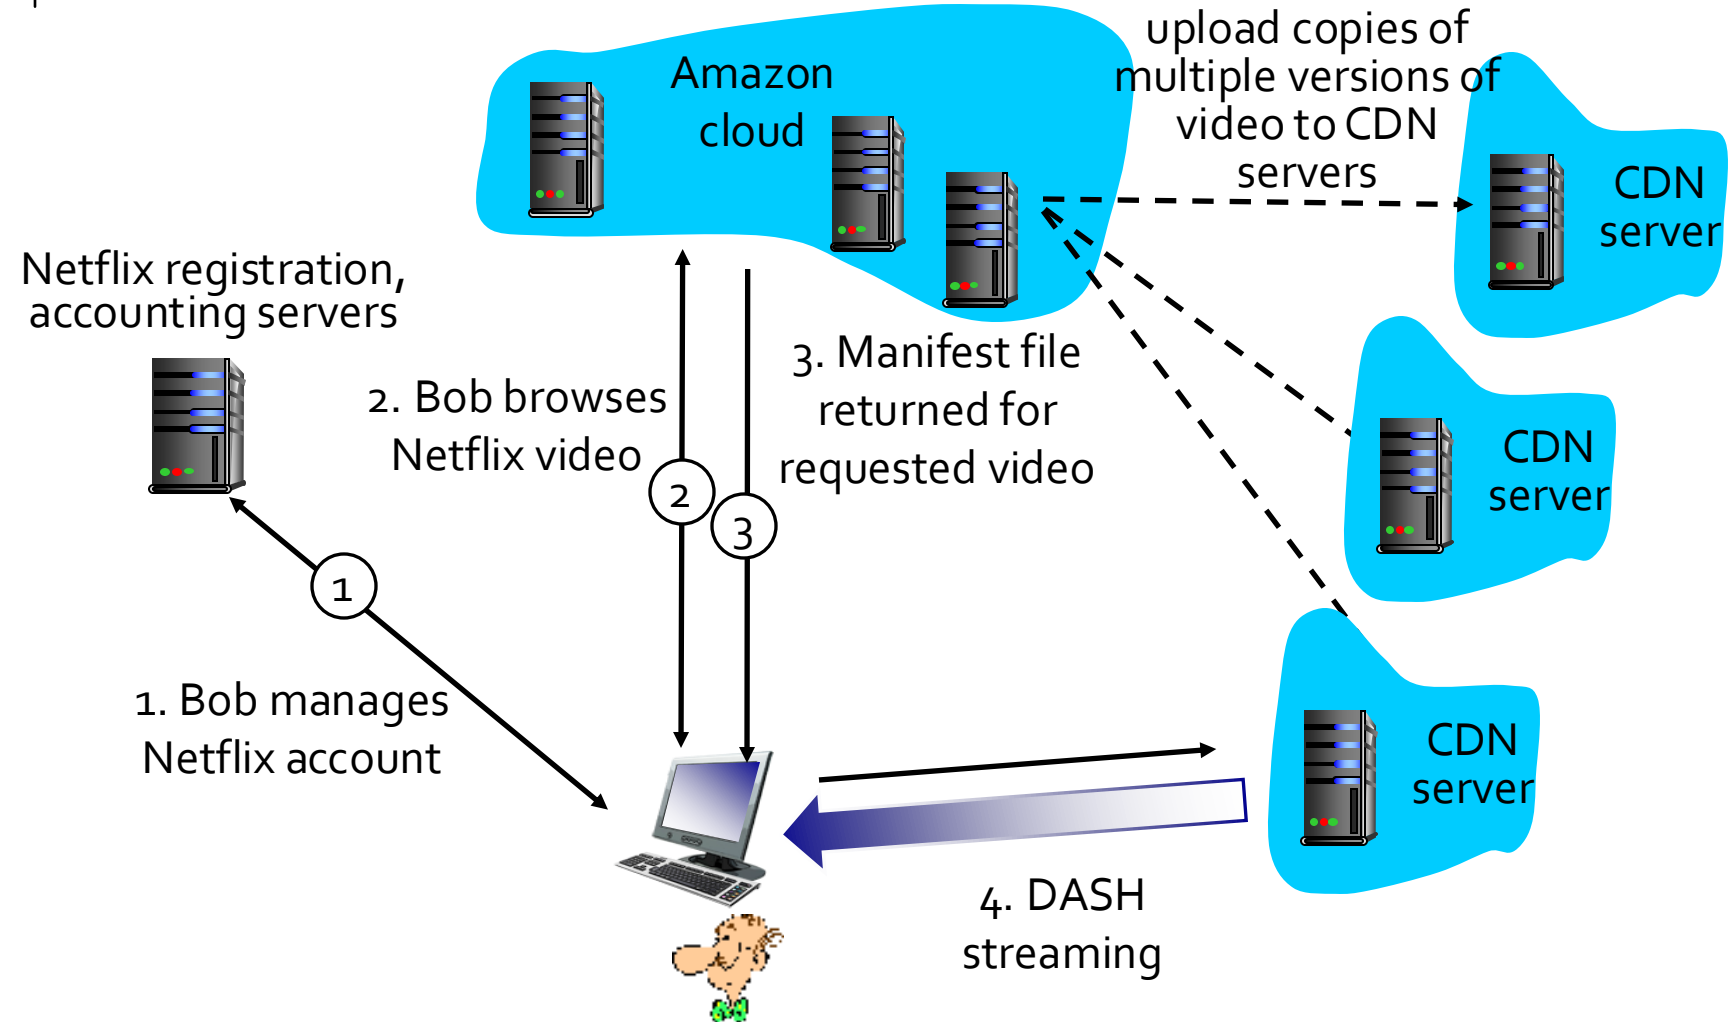
# CDN content access: a closer look

- Bob (client) requests video `http://netcinema.com/6Y7B23V`
  - video stored in CDN at `http://KingCDN.com/NetC6y&B23V`





# Netflix



# Summary

## Today:

- The eMail: SMTP, POP<sub>3</sub>
- P2P application architecture: BitTorrent
- Video streaming, CDNs

## Canvas discussion:

- Reflection
- Exit ticket

## Next time:

- read 3.1, 3.2, 3.3 and 3.4 of K&R (Transport layer)
- follow on Canvas! material and announcements

# Any questions?