Reliable data transfer - Noisy channel

CE 352, Computer Networks

Salem Al-Agtash

Lecture 9

Recap

- Transport protocols:
  - Reliable channel – Simple, rdt1.0
  - Unreliable channel – Stop-and-Wait
    - rdt2.0: channel with bit errors
    - rdt2.1: distorted ACK/NAK)
    - rdt2.2: NAK – free protocol

Today

    - rdt3.0: channel with errors and loss
  - Unreliable channel - pipelined protocols
    - Go-Back-N
    - Selective repeat
  - TCP protocol

rdt1.0: reliable transfer over a reliable channel

- underlying channel perfectly reliable
  - no bit errors
  - no loss of packets
- separate FSMs for sender, receiver:
  - sender sends data into underlying channel
  - receiver reads data from underlying channel

Simple
provides neither
flow nor error control

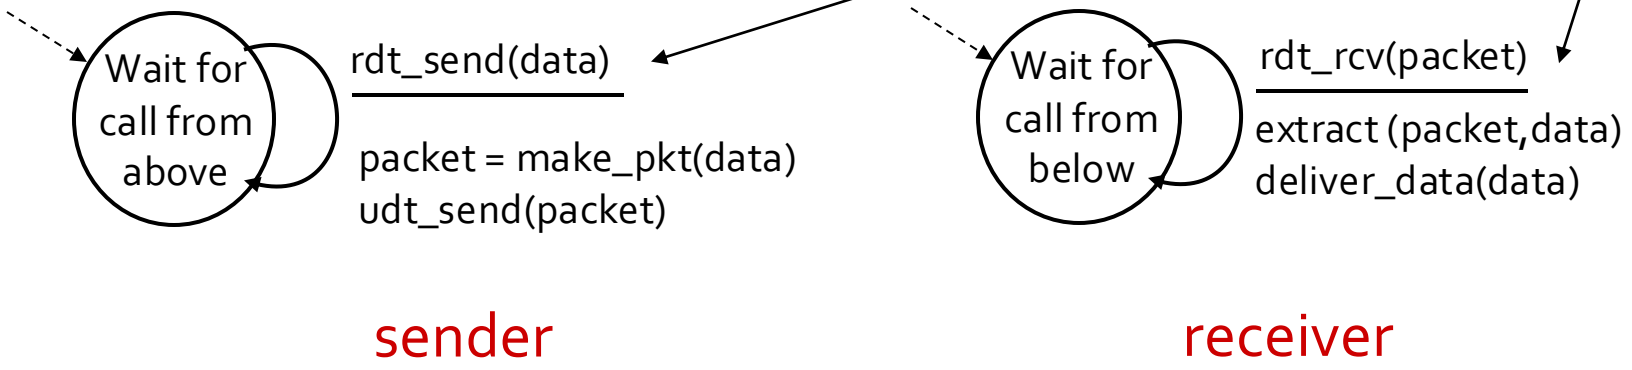

rdt1.0 in action

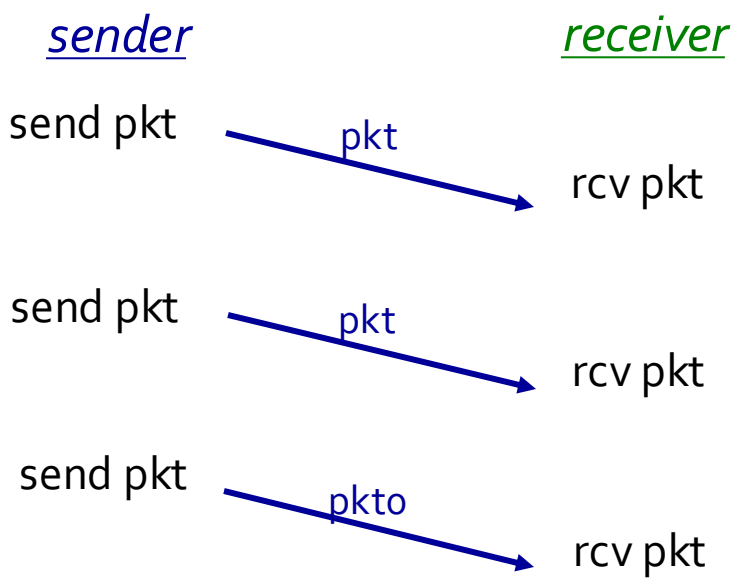

No error and no loss (reliable communication channel)

rdt2.0: channel with bit errors (unreliable channel)

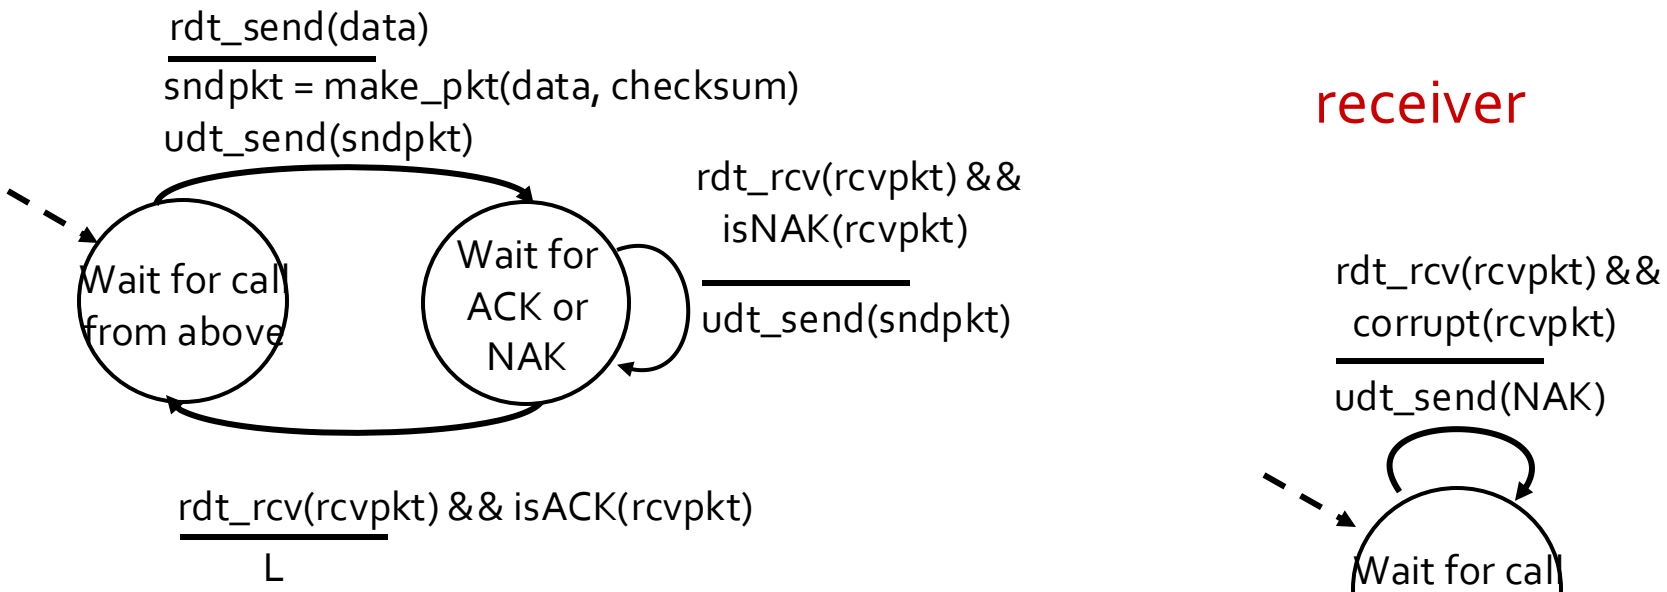

stop and wait

rdt2.0 in action

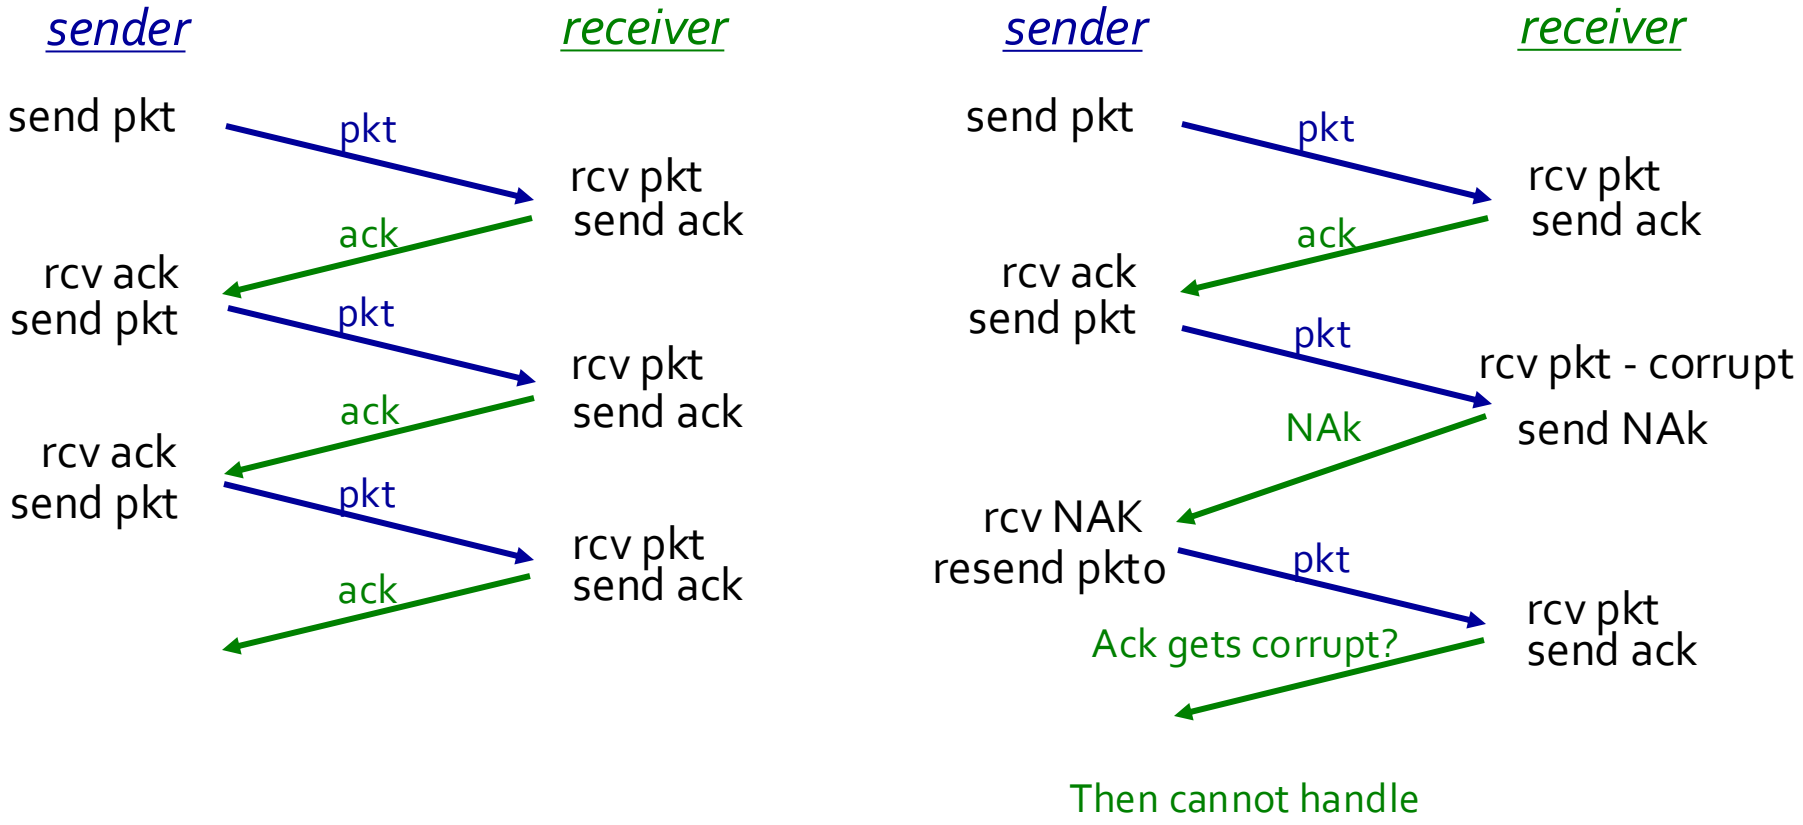

Possible error in Ack but still assuming no loss of packets

rdt2.1: sender handles distorted ACK/NAKs

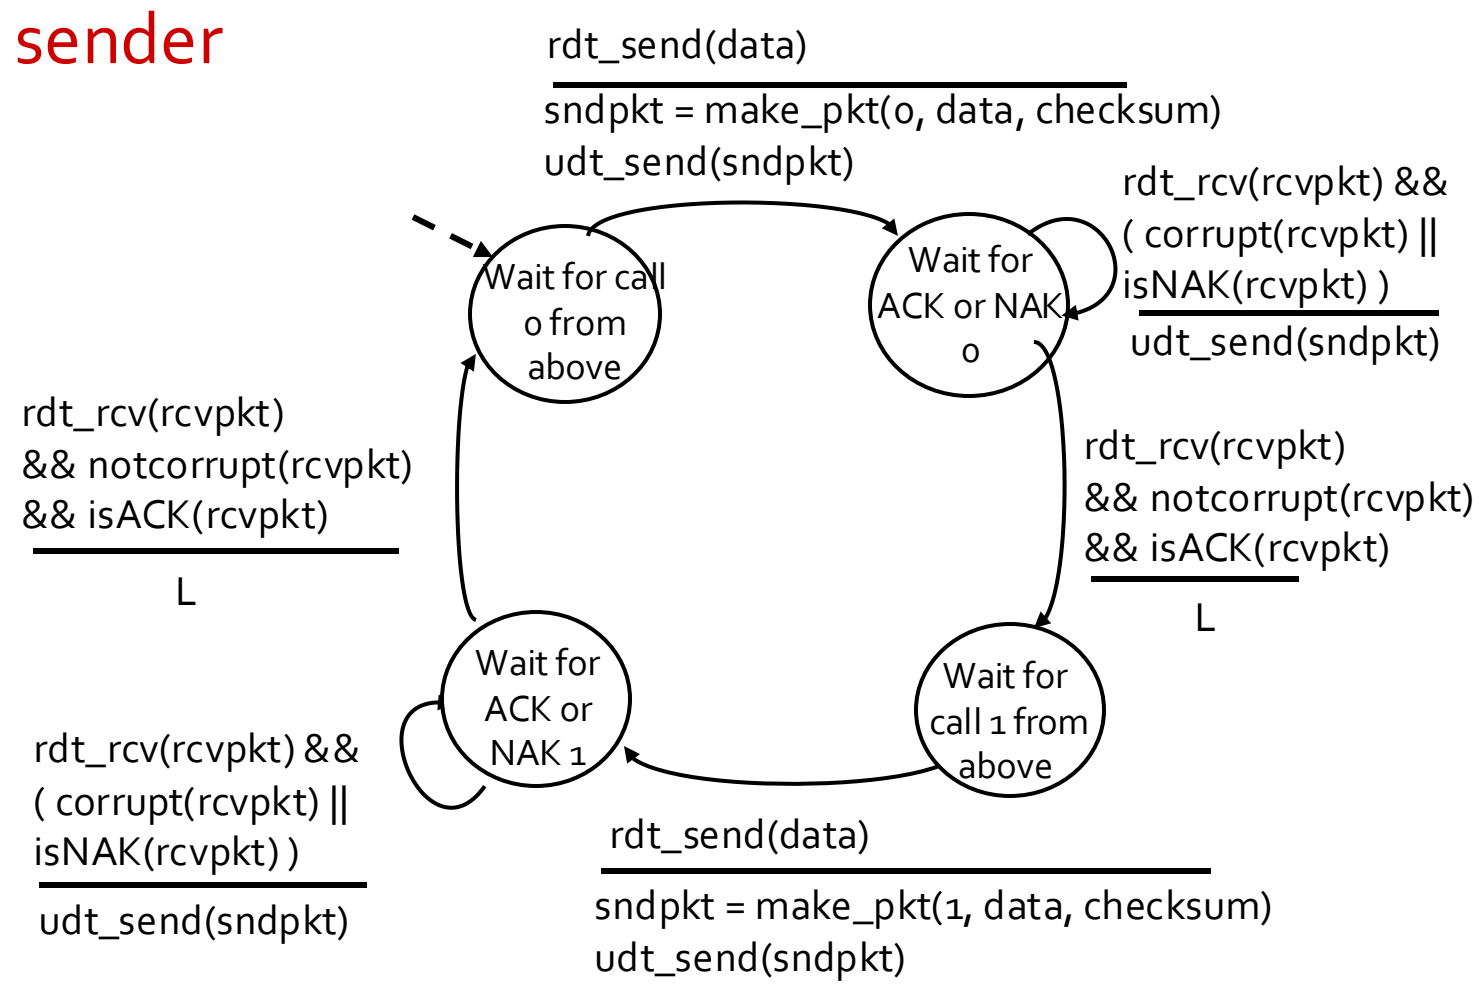

rdt2.1: receiver handles distorted ACK/NAKs

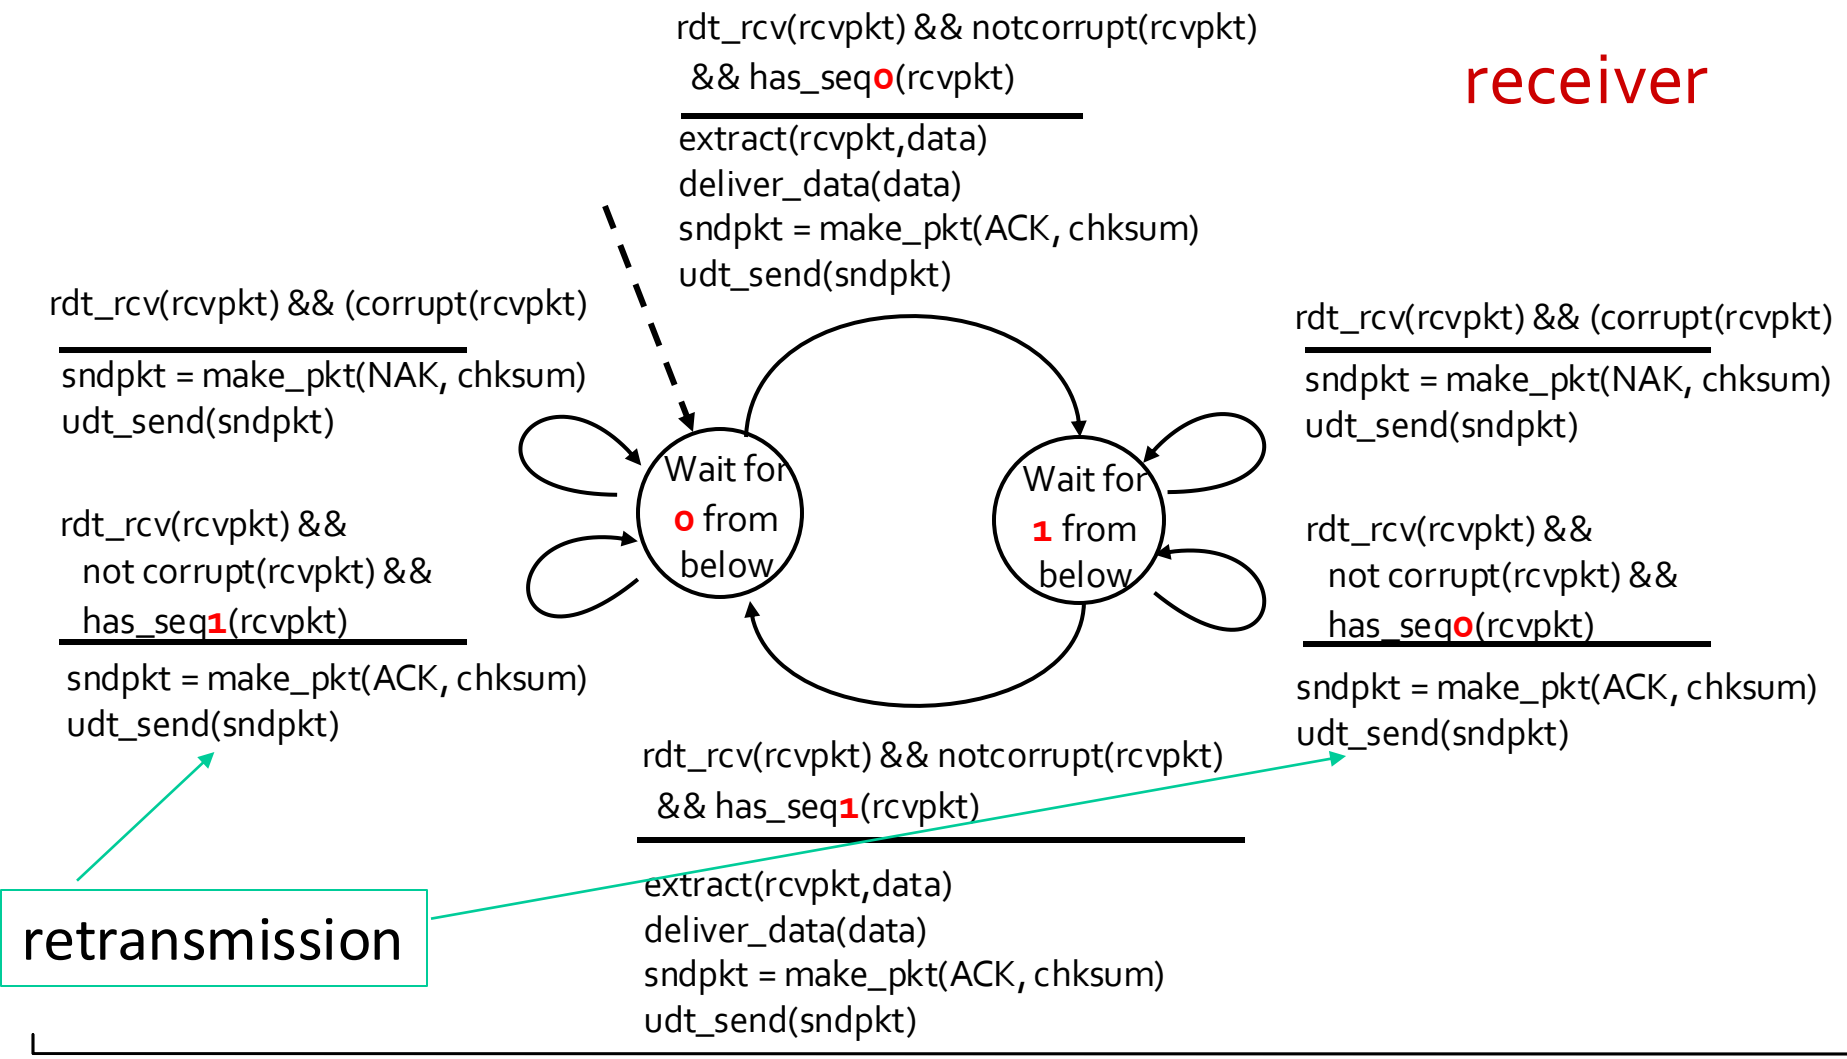

rdt2.1 in action

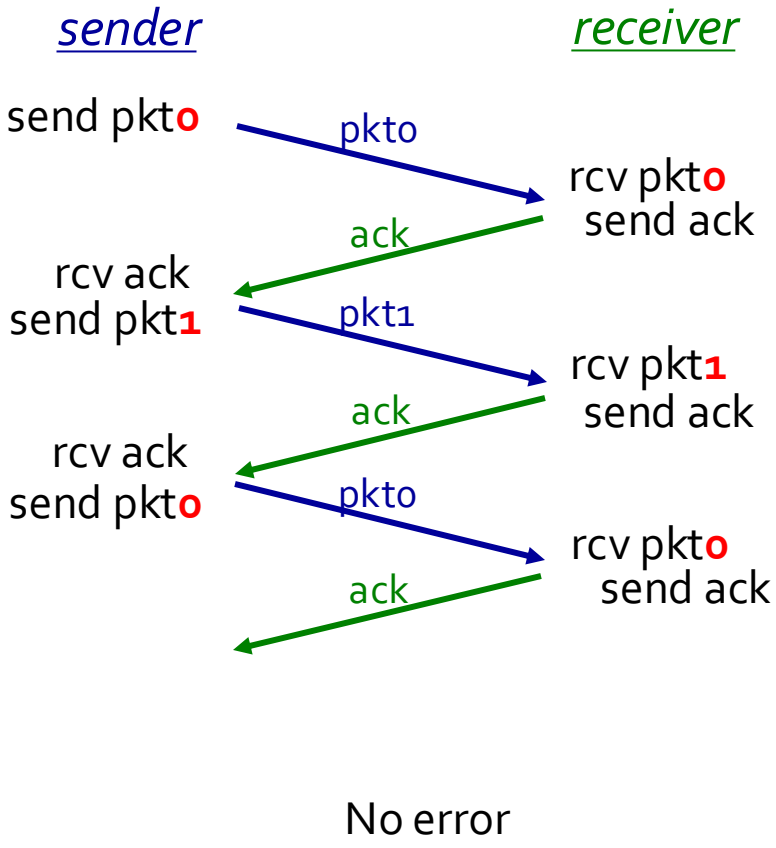

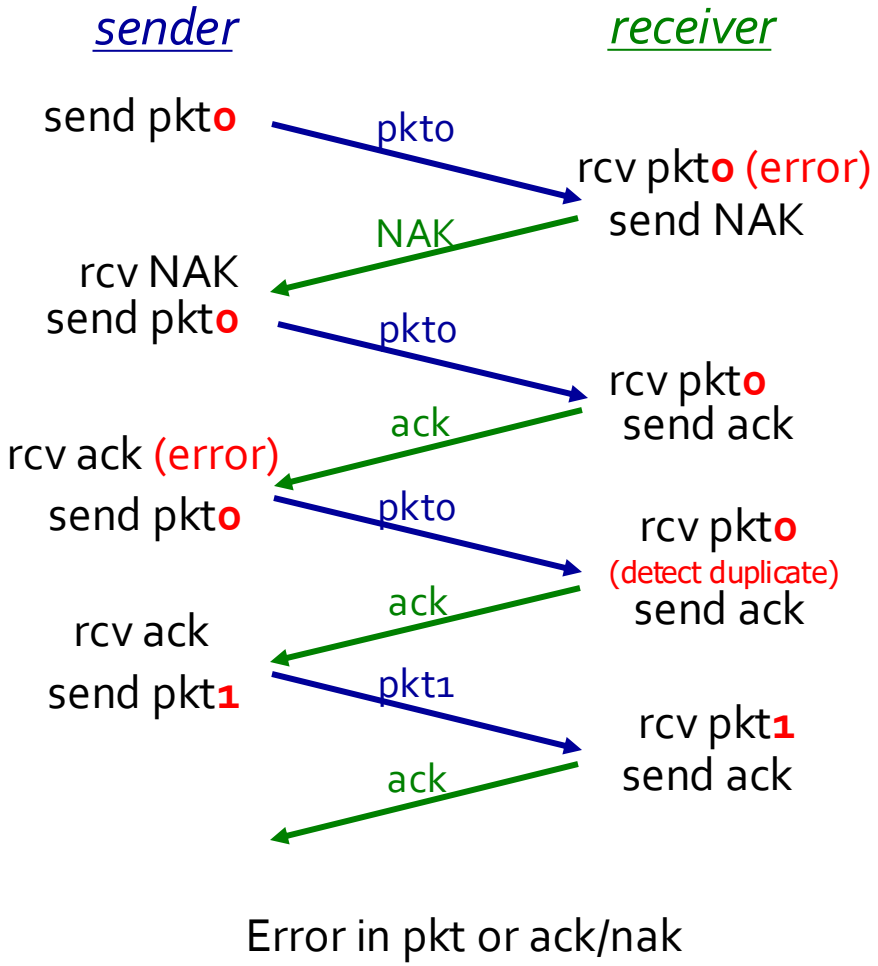

rdt2.2: a NAK-free protocol

- same functionality as rdt2.1, using ACKs only
- instead of NAK, receiver sends ACK for last pkt received OK
  - receiver must *explicitly* include seq # of pkt being ACKed
- duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

rdt2.2: NAK-free sender

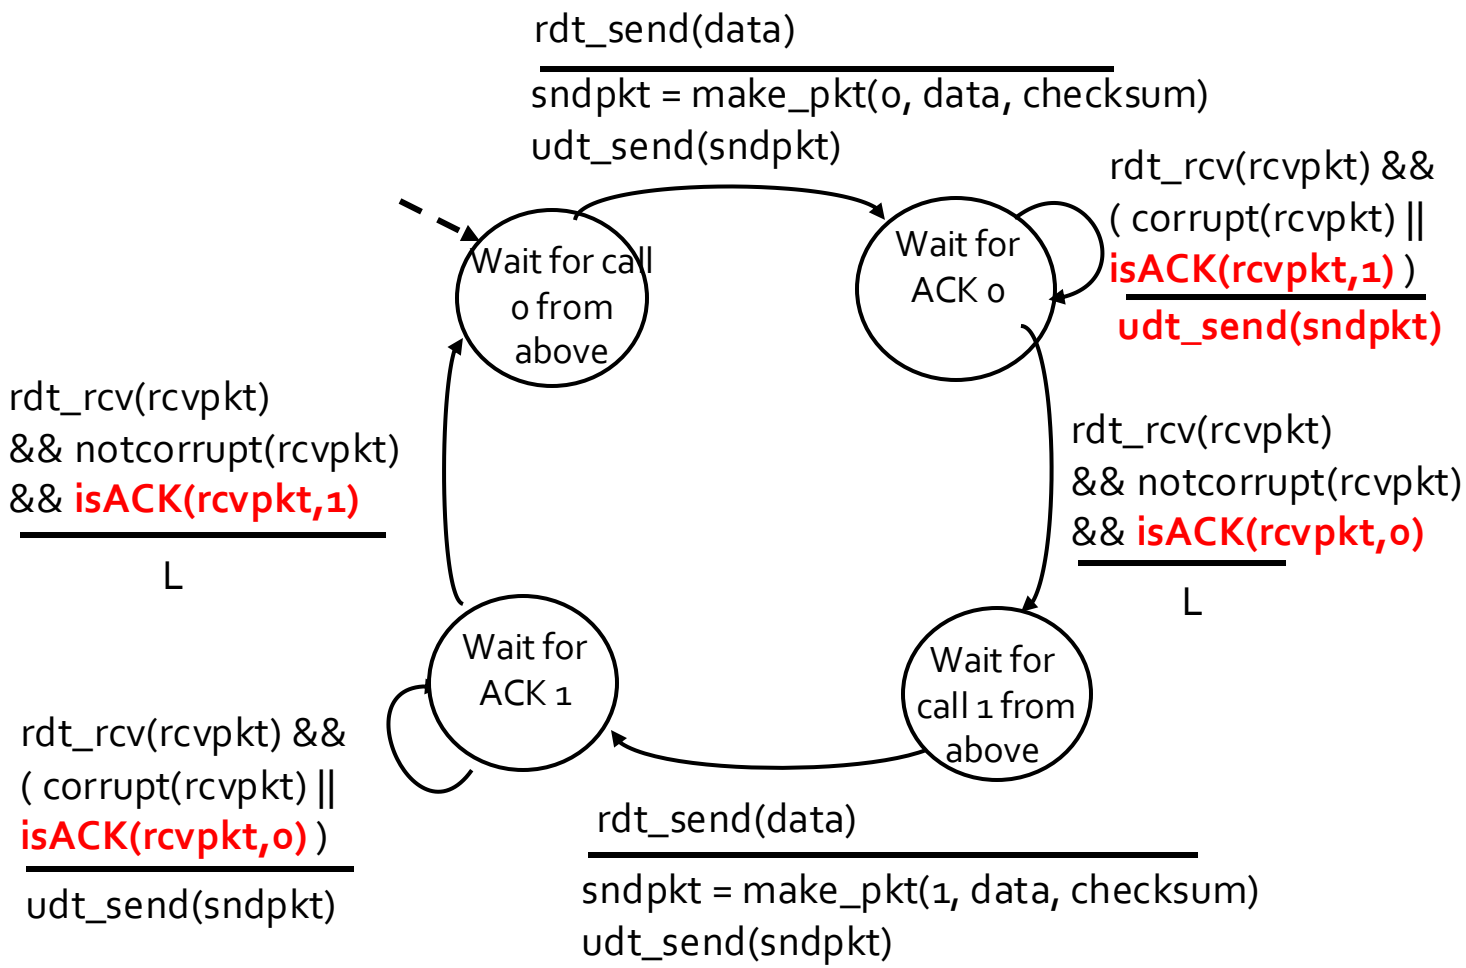

rdt2.2: NAK-free receiver

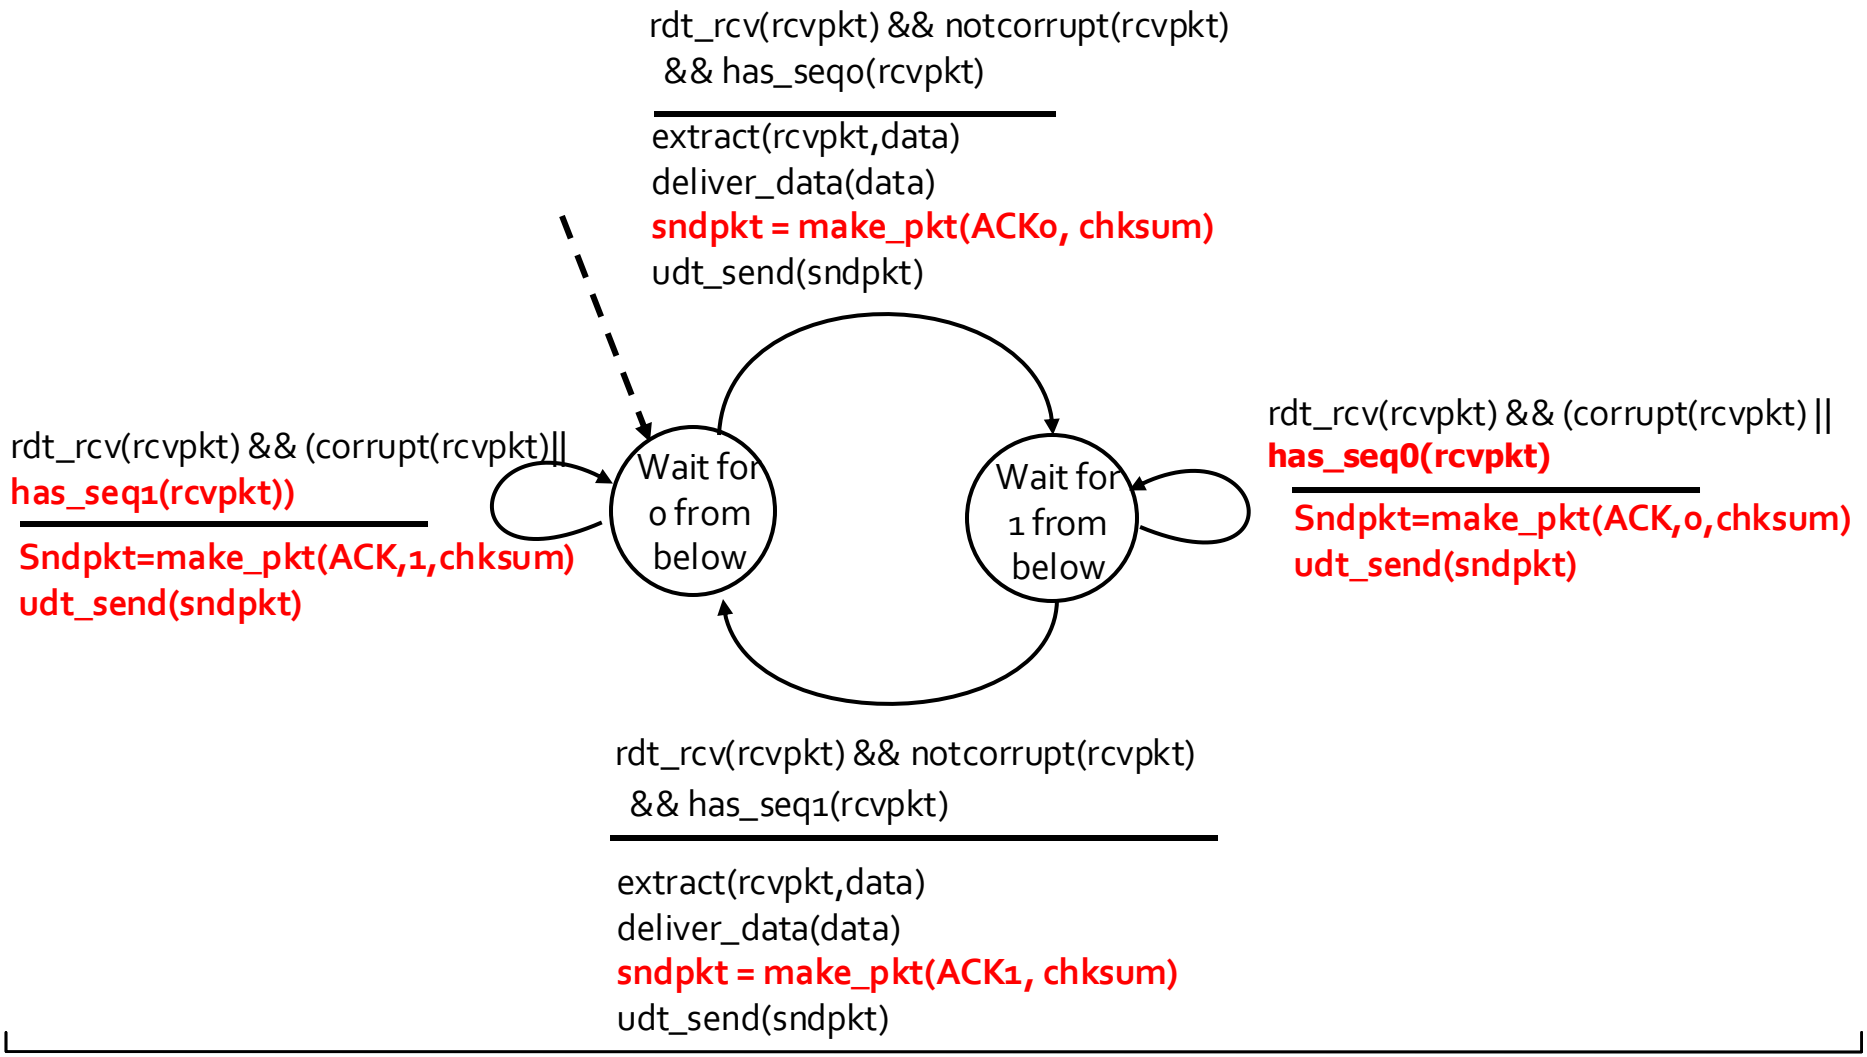

rdt2.2 in action

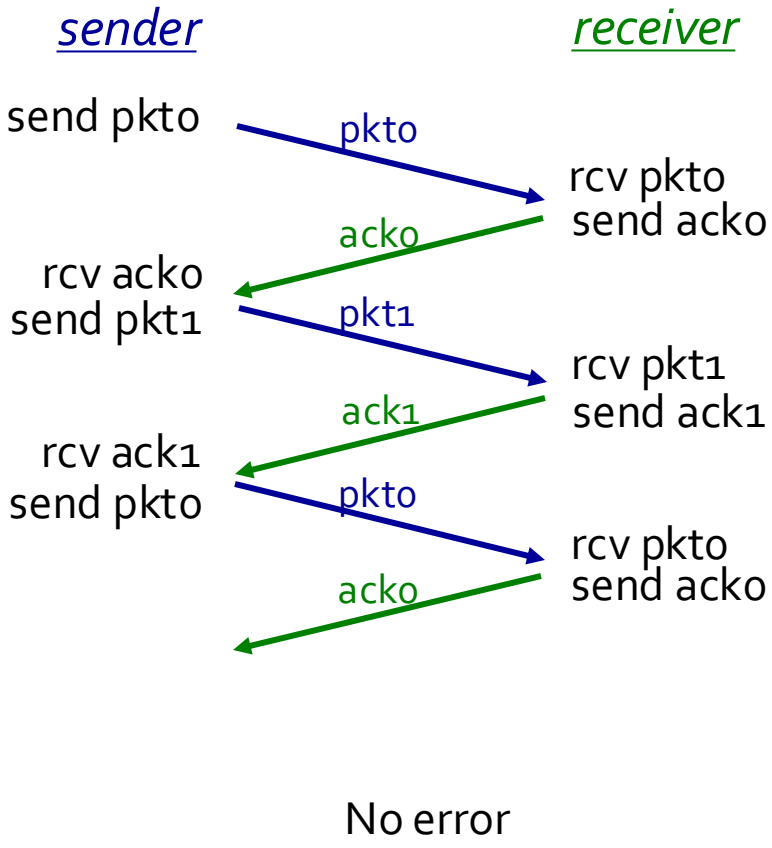

No error

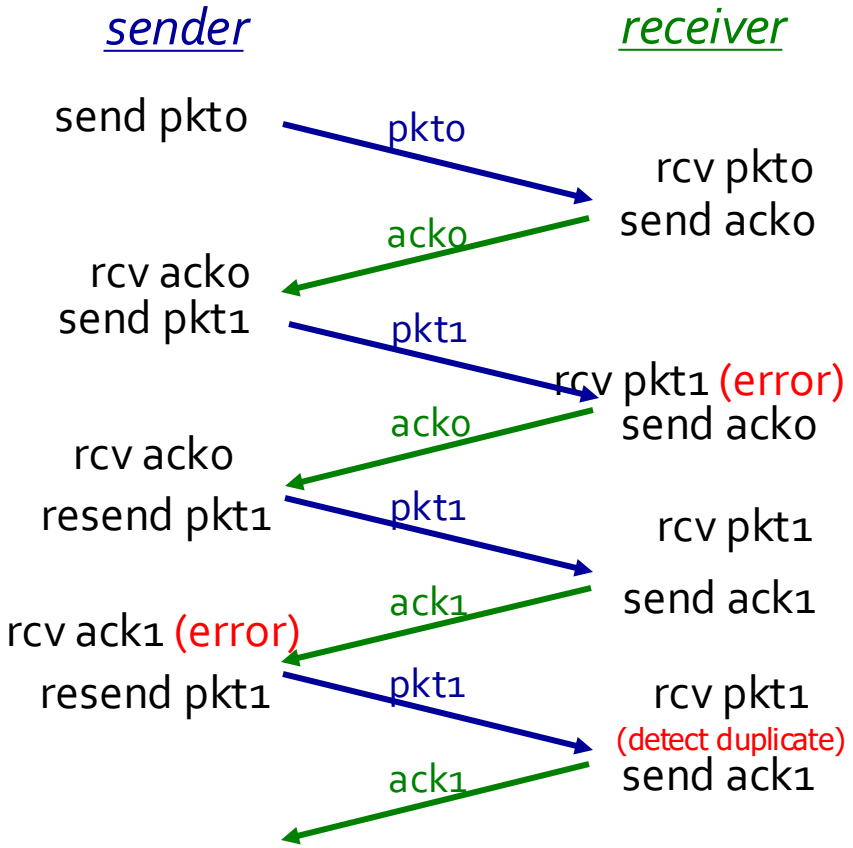

Error in pkt or ack/nak

rdt3.0: channels with errors *and* loss

New assumption:

- underlying channel can also lose packets (data, ACKs)
  - checksum, seq. #, ACKs, retransmissions will be of help … but not enough

Approach:

- sender waits "reasonable" amount of time for ACK
- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
  - retransmission will be duplicate, but seq. #'s already handles this
  - receiver must specify seq # of pkt being ACKed
- requires countdown timer

rdt3.0 sender

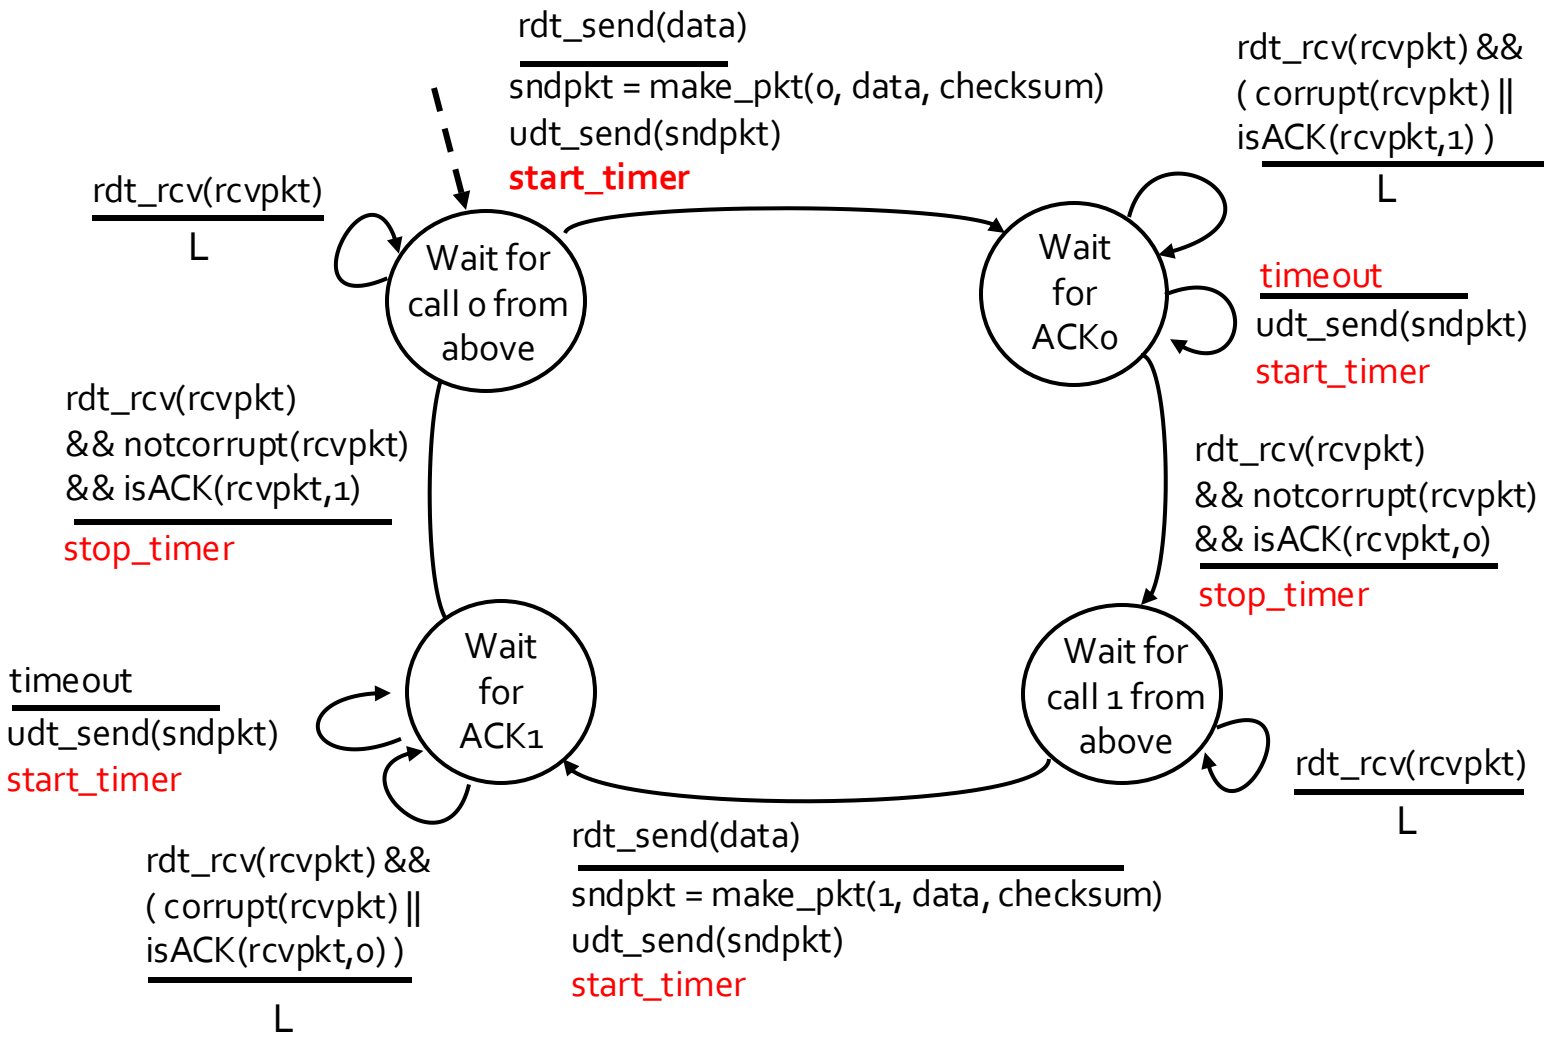

rdt3.0 receiver (rdt2.2 receiver)

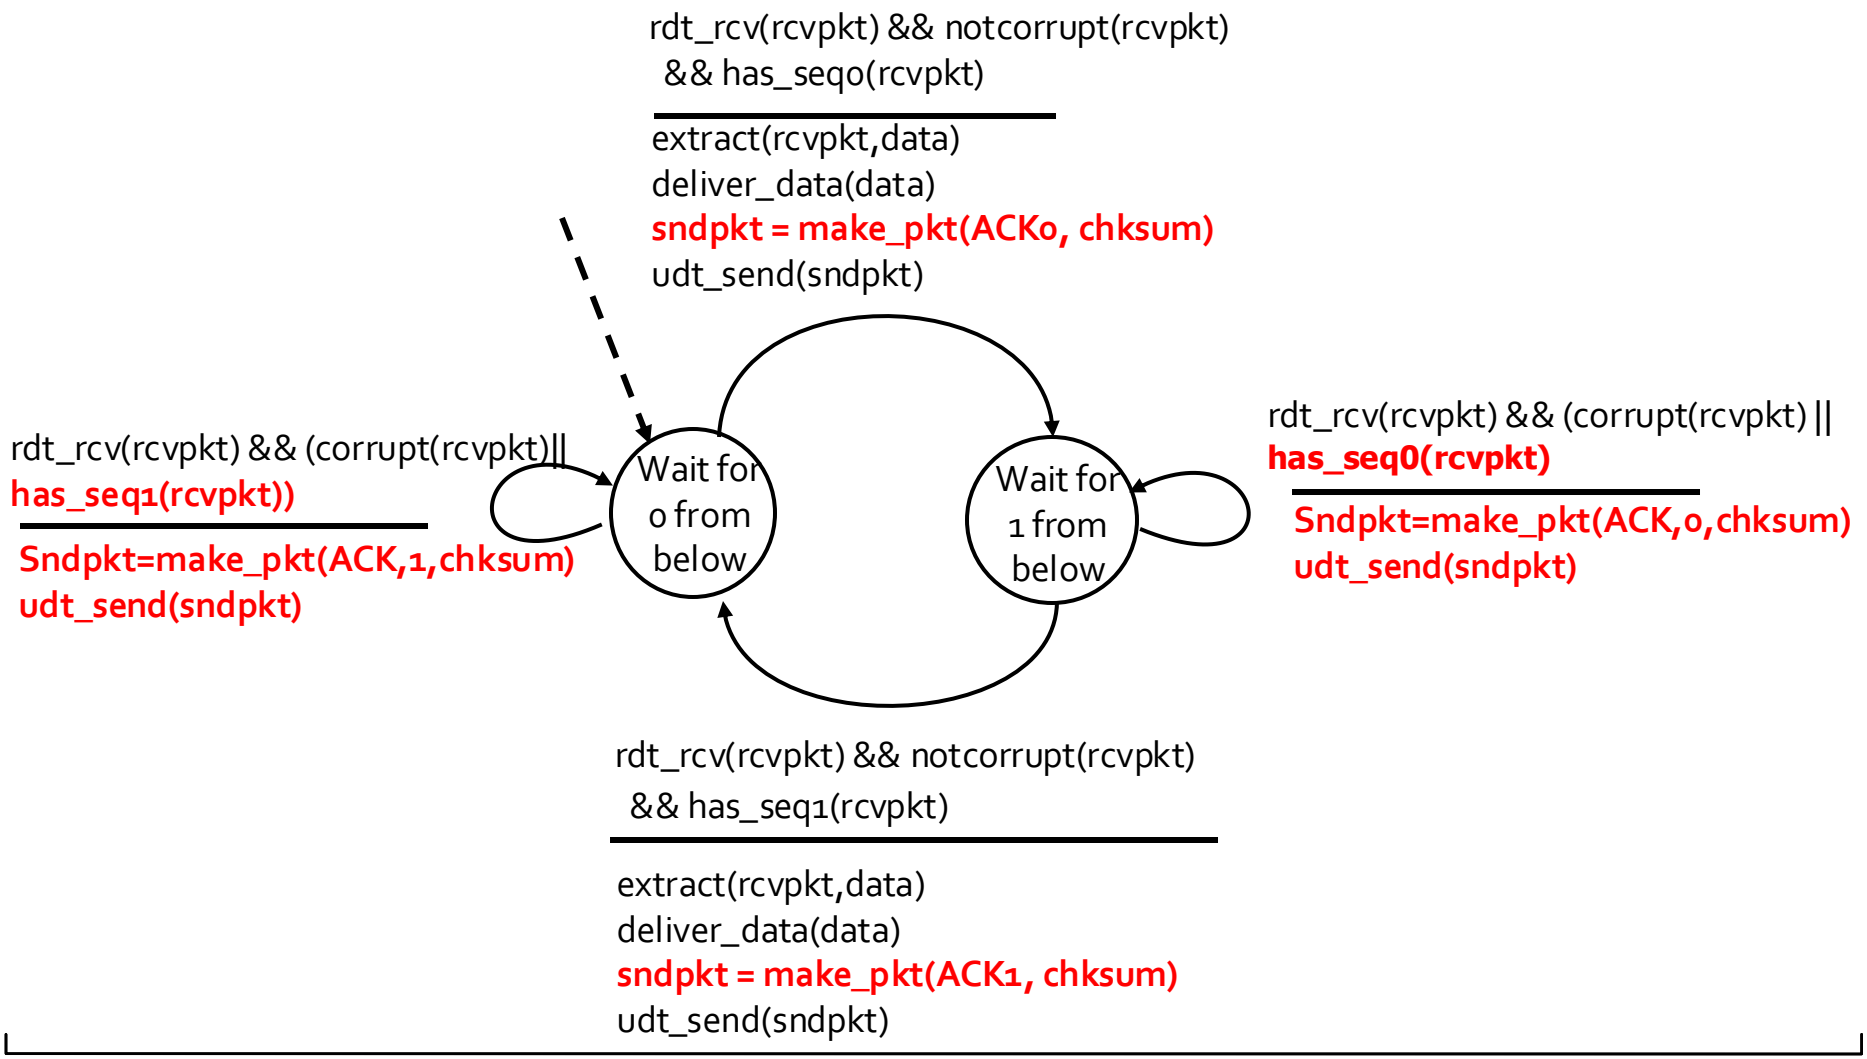

rdt3.0 in action

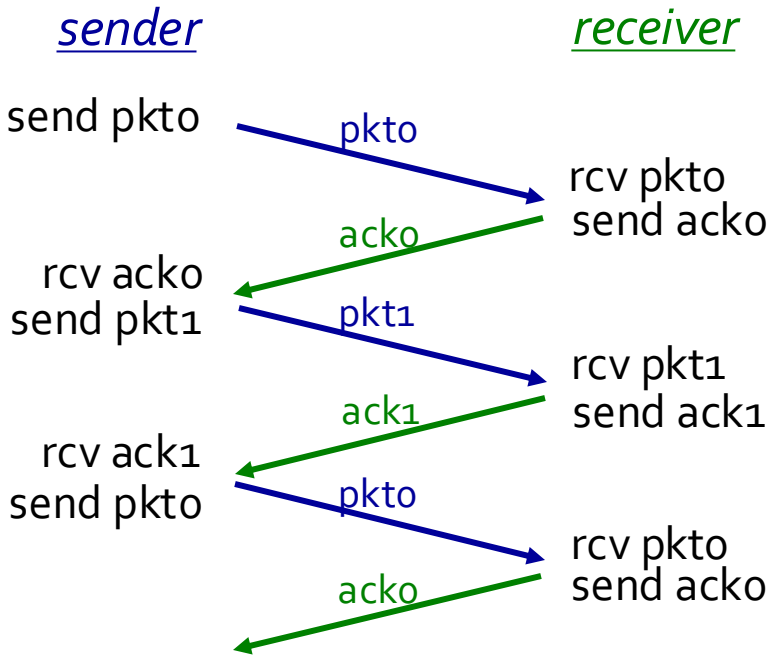

(a) no loss

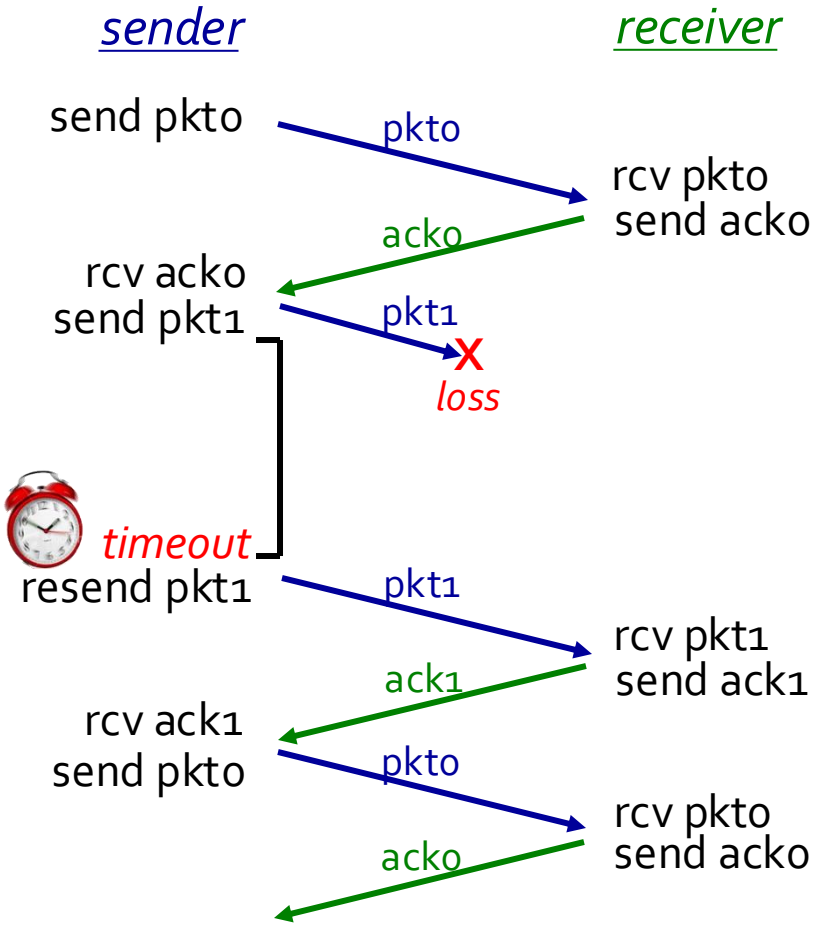

(b) packet loss

rdt3.0 in action

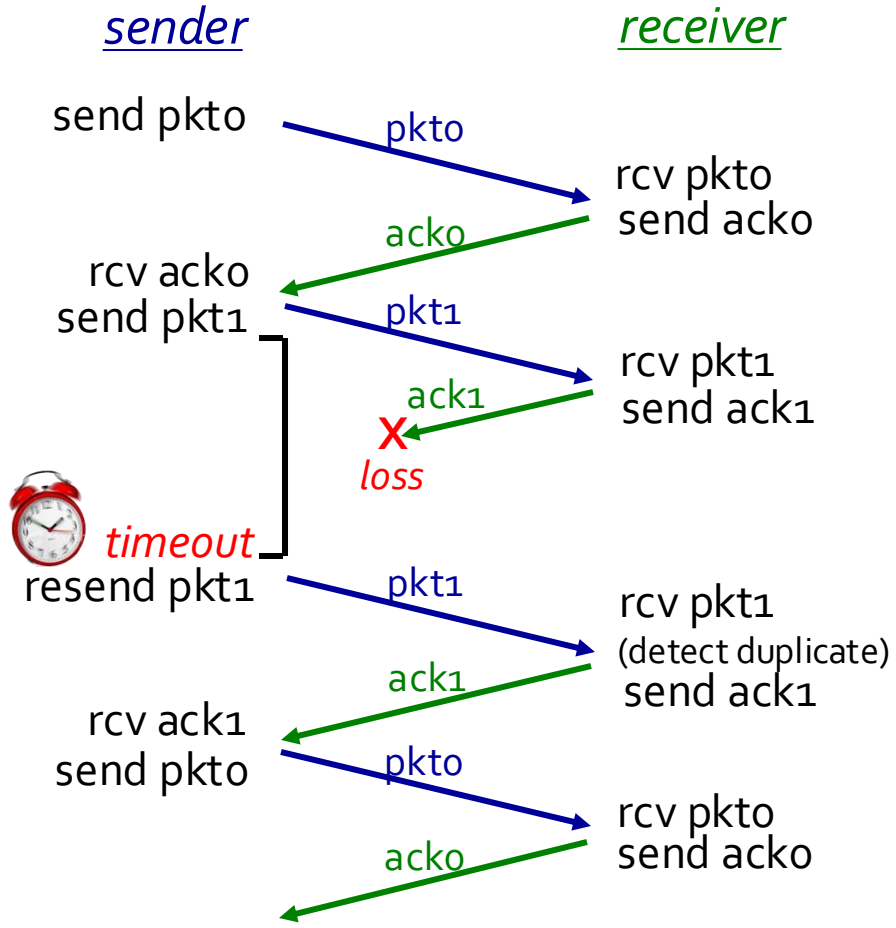

(c) ACK loss

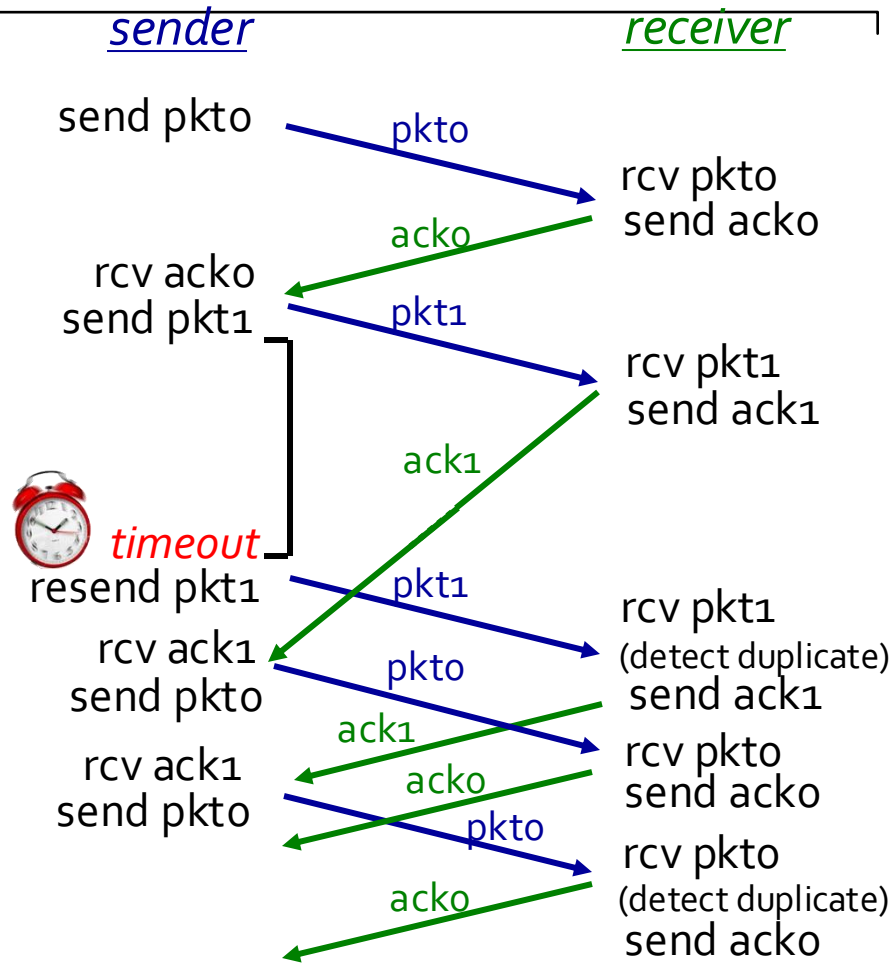

(d) premature timeout/ delayed ACK

rdt3.0: stop-and-wait operation

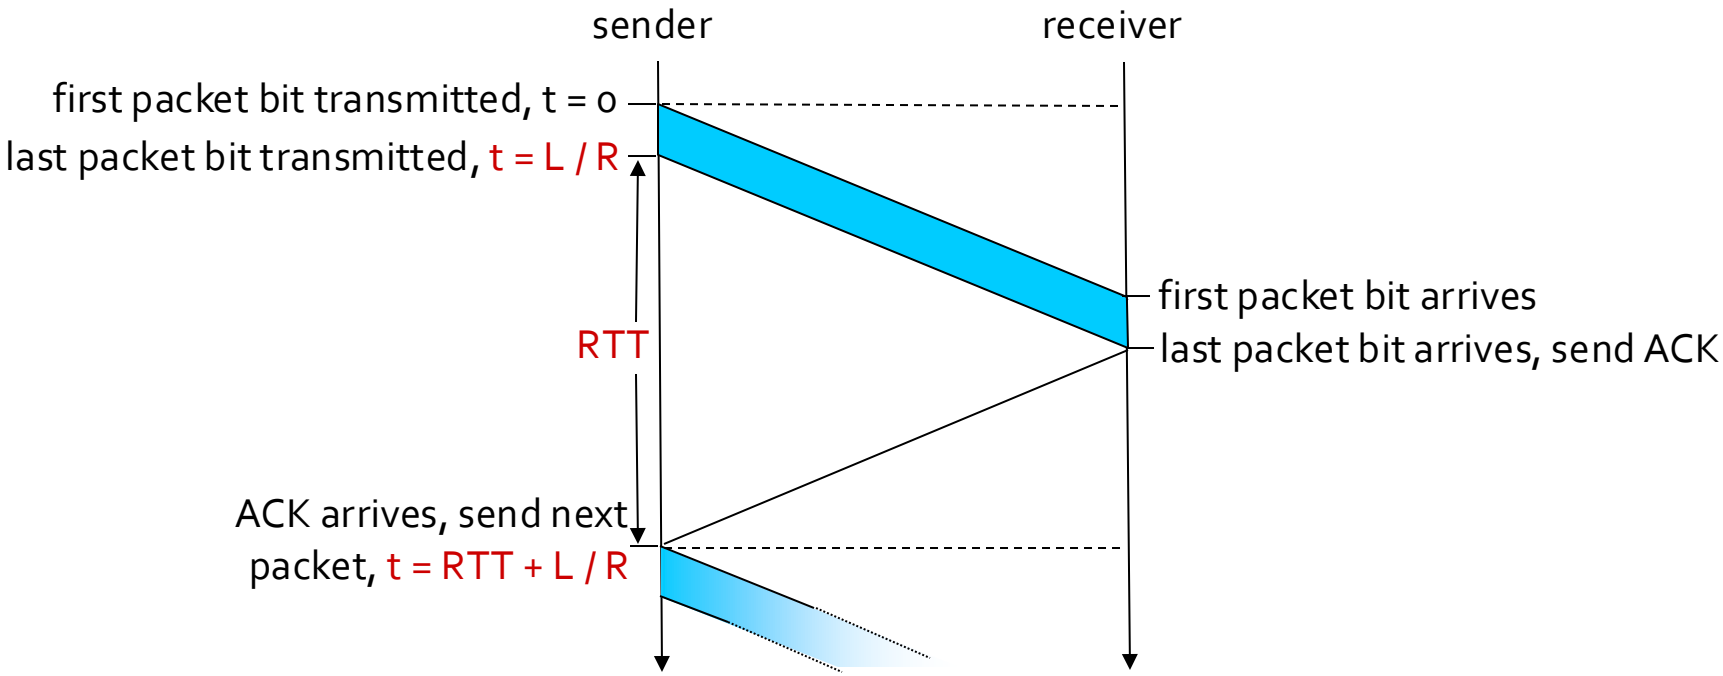

$$U_{sender} = \frac{L / R}{RTT + L / R}$$

Performance of rdt3.0

- rdt3.0 is correct, but performance is not good
- e.g.: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

$$D_{trans} = \frac{L}{R} = \frac{8000\ bits}{10^9\ bits/sec} = 8\ microsecs$$

  - U_{sender}: *utilization* – fraction of time sender busy sending

$$U_{sender} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$

  - if RTT=30 msec, 1KB pkt every 30 msec: 33kB/sec (267kbps) thruput over 1 Gbps link
- network protocol limits use of physical resources!

Pipelined protocols

- pipelining: sender allows multiple, "in-flight", yet-to-be-acknowledged pkts
  - range of sequence numbers must be increased
  - buffering at sender and/or receiver
- two generic forms of pipelined protocols: *go-Back-N, selective repeat*

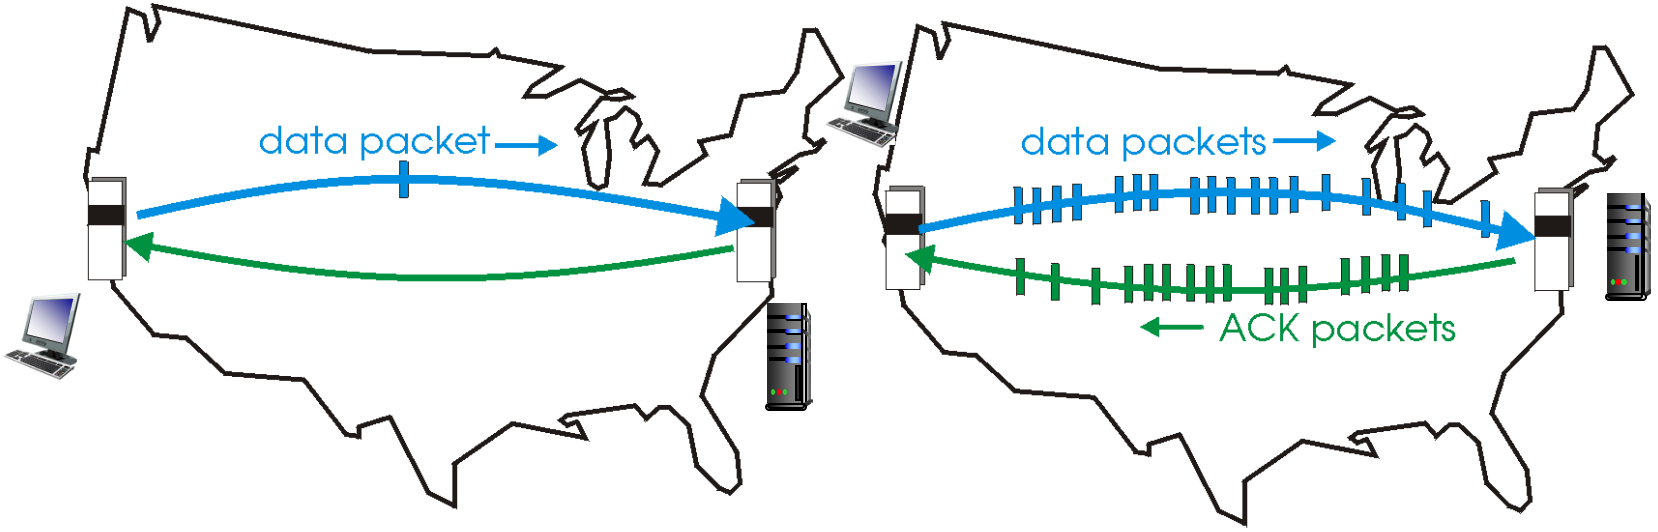

(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

Pipelining: increased utilization

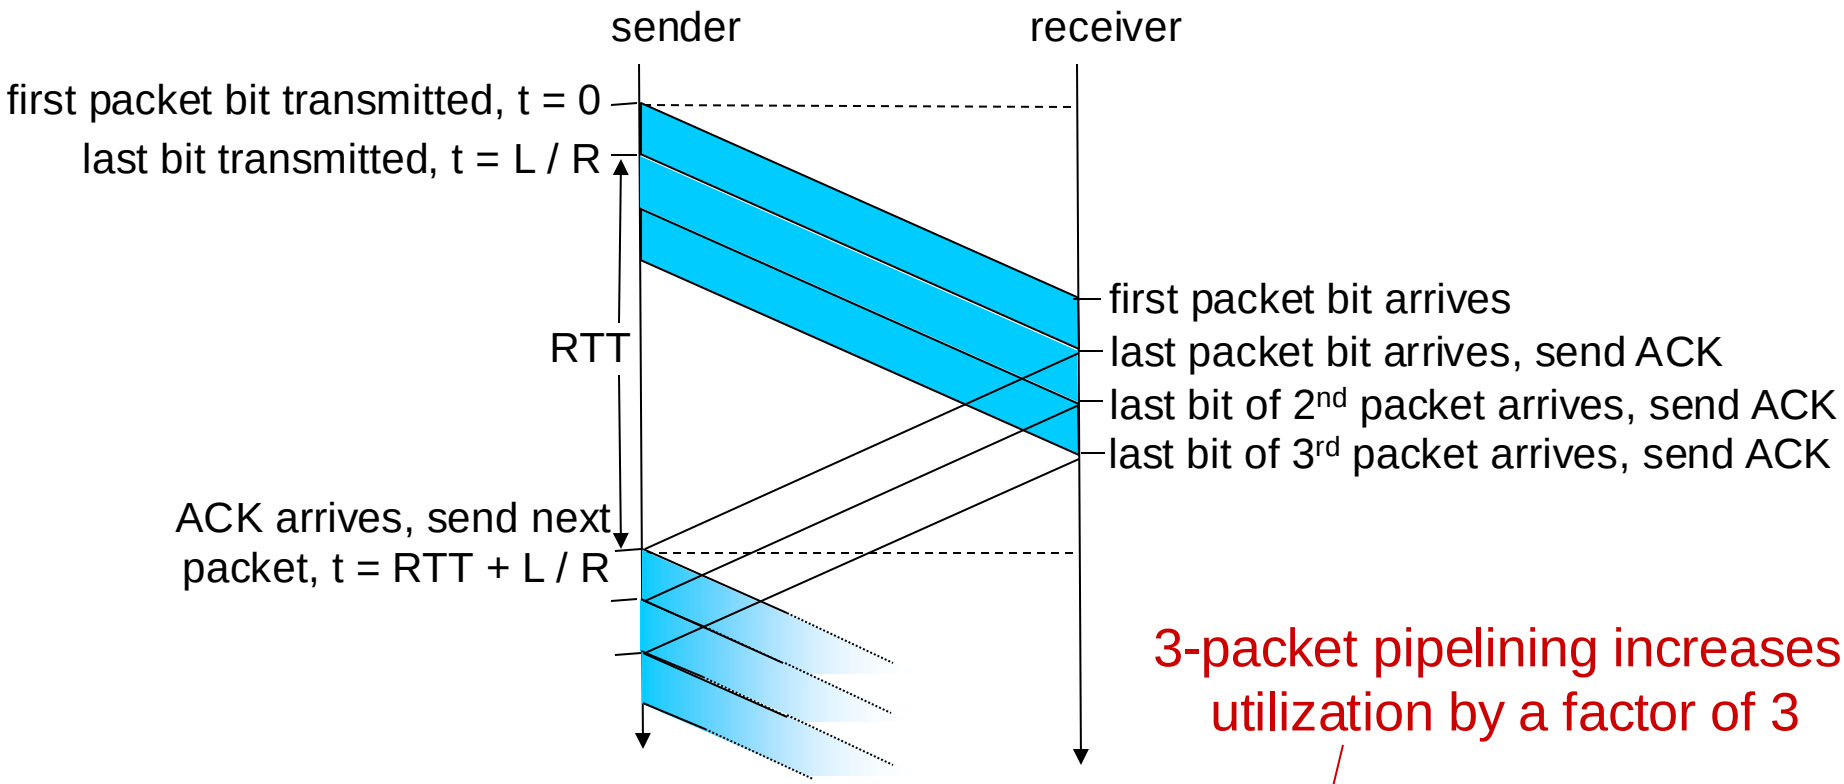

3-packet pipelining increases utilization by a factor of 3

$$U_{sender} = \frac{3L / R}{RTT + L / R} = \frac{.0024}{30.008} = 0.00081$$

Pipelined protocols: overview

Go-back-N:

- sender can have up to N unacked packets in pipeline
- receiver only sends *cumulative ack*
  - doesn't ack packet if there's a gap
- sender has timer for oldest unacked packet
  - when timer expires, retransmit *all* unacked packets

Selective Repeat:

- sender can have up to N unack'ed packets in pipeline
- receiver sends *individual ack* for each packet
- sender maintains timer for each unacked packet
  - when timer expires, retransmit only that unacked packet

Go-Back-N: sender

- k-bit seq # in pkt header
- "window" of up to N, consecutive unack'ed pkts allowed
- ACK(n): ACKs all pkts up to, including seq # n - *"cumulative ACK"*
  - may receive duplicate ACKs (see receiver)
- timer for oldest in-flight pkt
- *timeout(n):* retransmit packet n and all higher seq # pkts in window

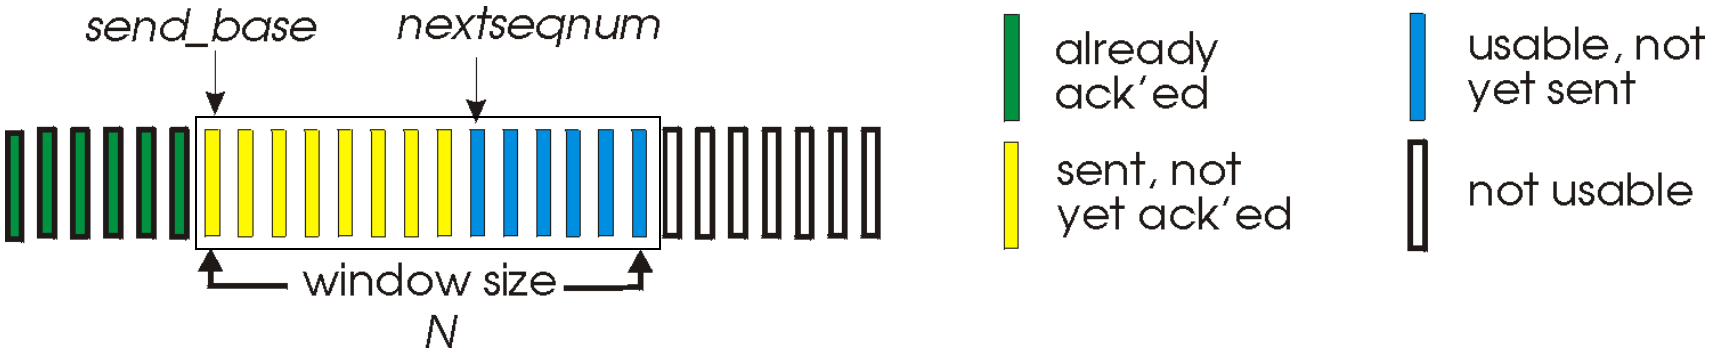

GBN in action

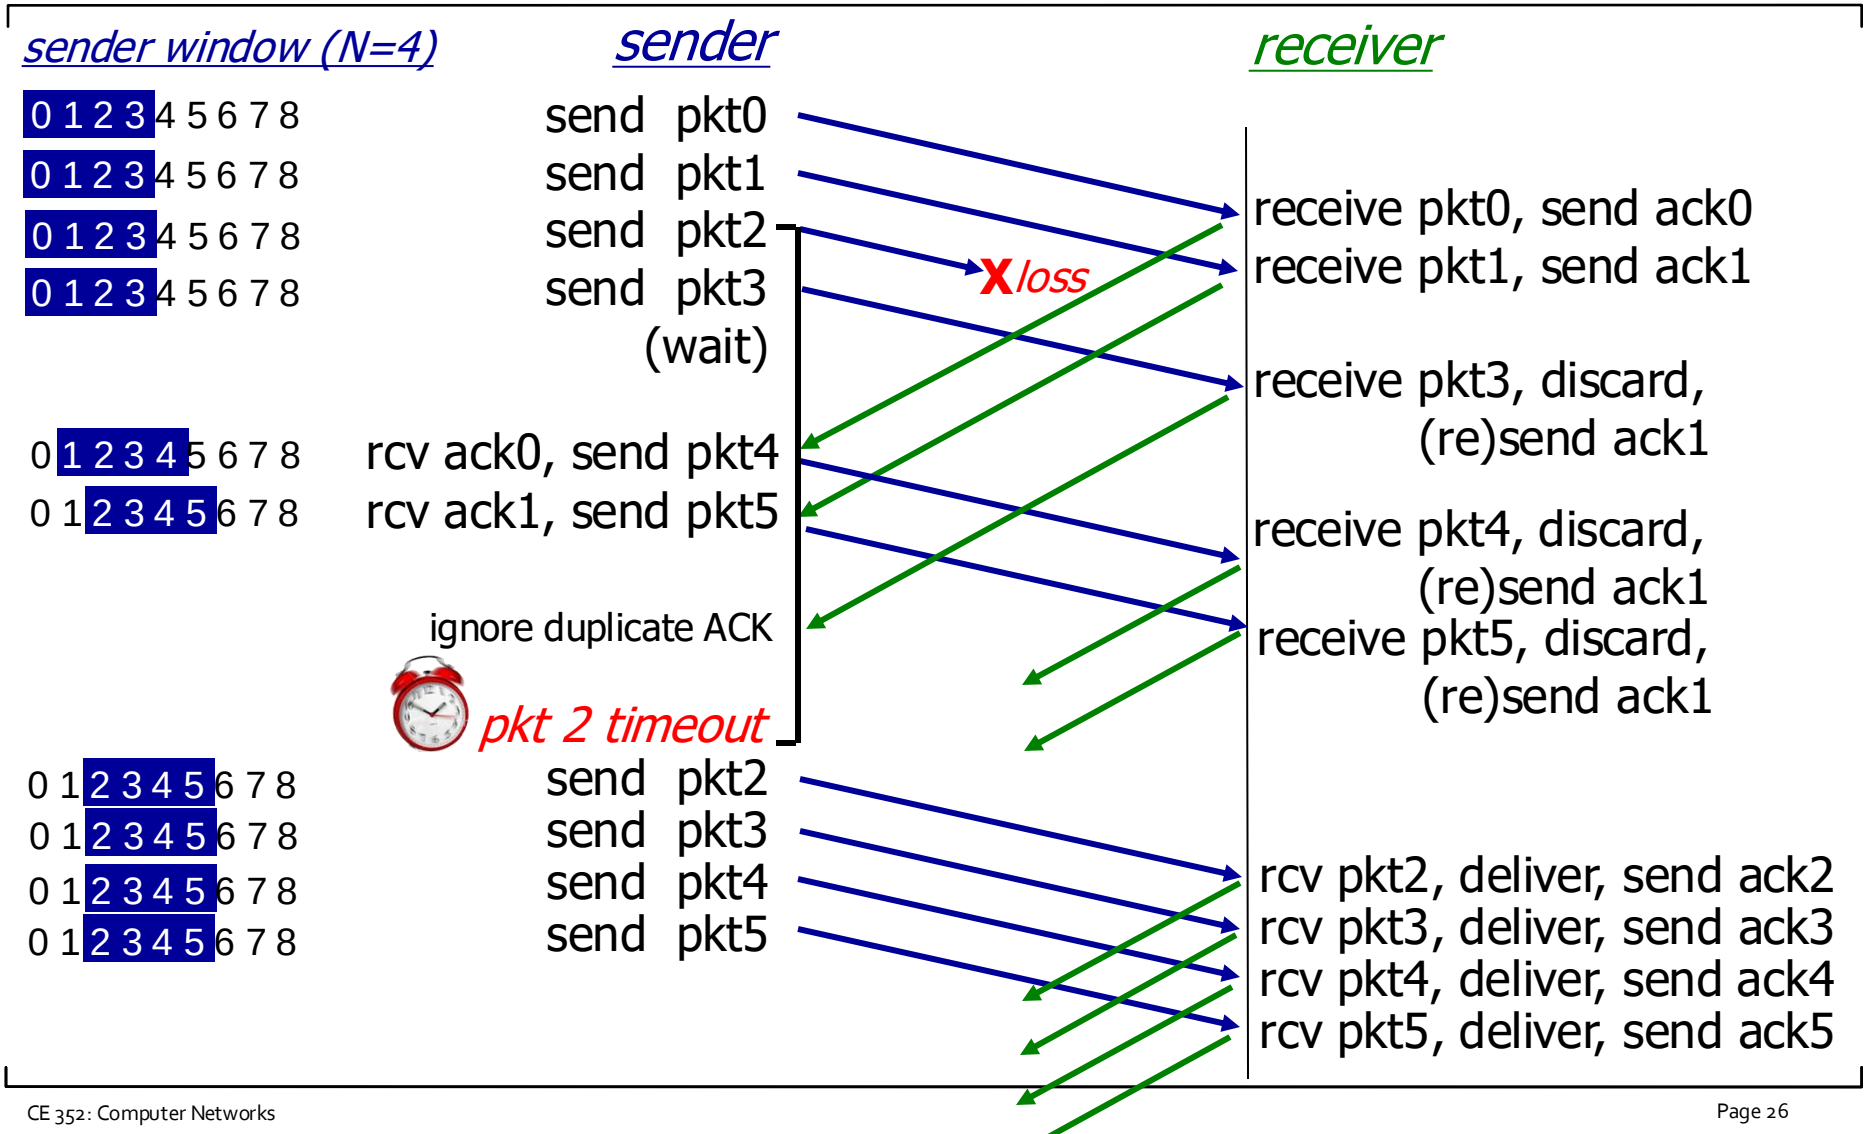

Selective repeat

- receiver *individually* acknowledges all correctly received pkts
  - buffers pkts, as needed, for eventual in-order delivery to upper layer
- sender only resends pkts for which ACK not received
  - sender timer for each unACKed pkt
- sender window
  - *N* consecutive seq #'s
  - limits seq #s of sent, unACKed pkts

Selective repeat: sender, receiver windows

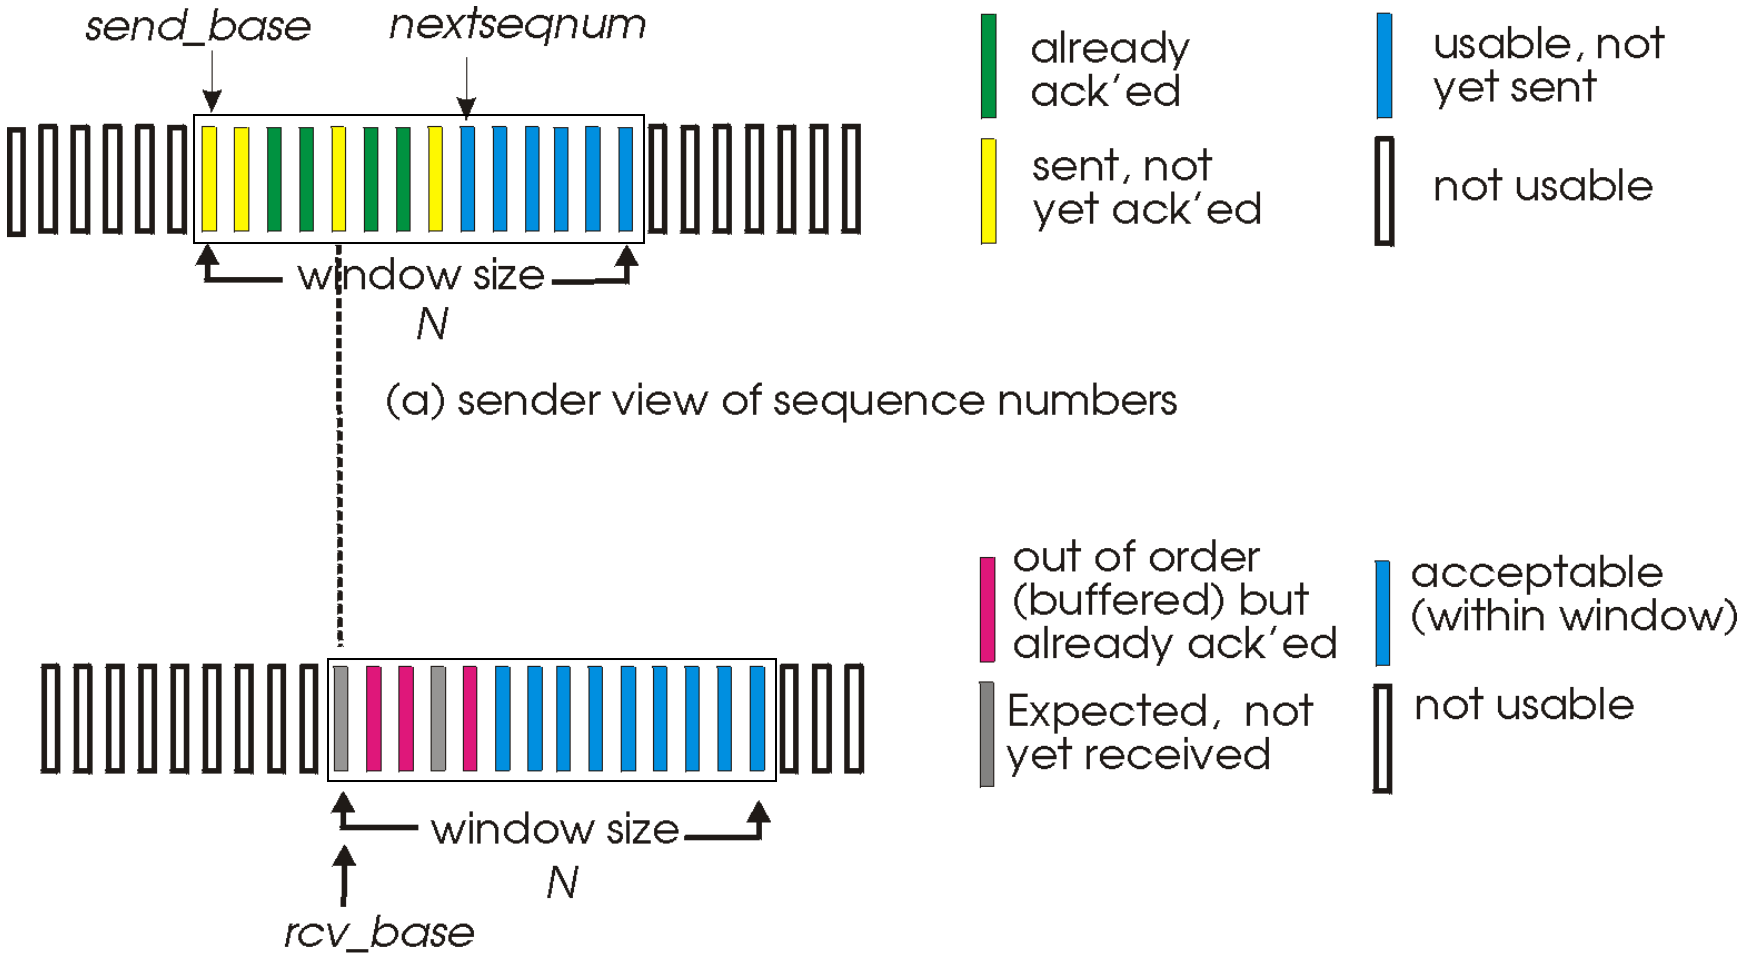

Selective repeat in action

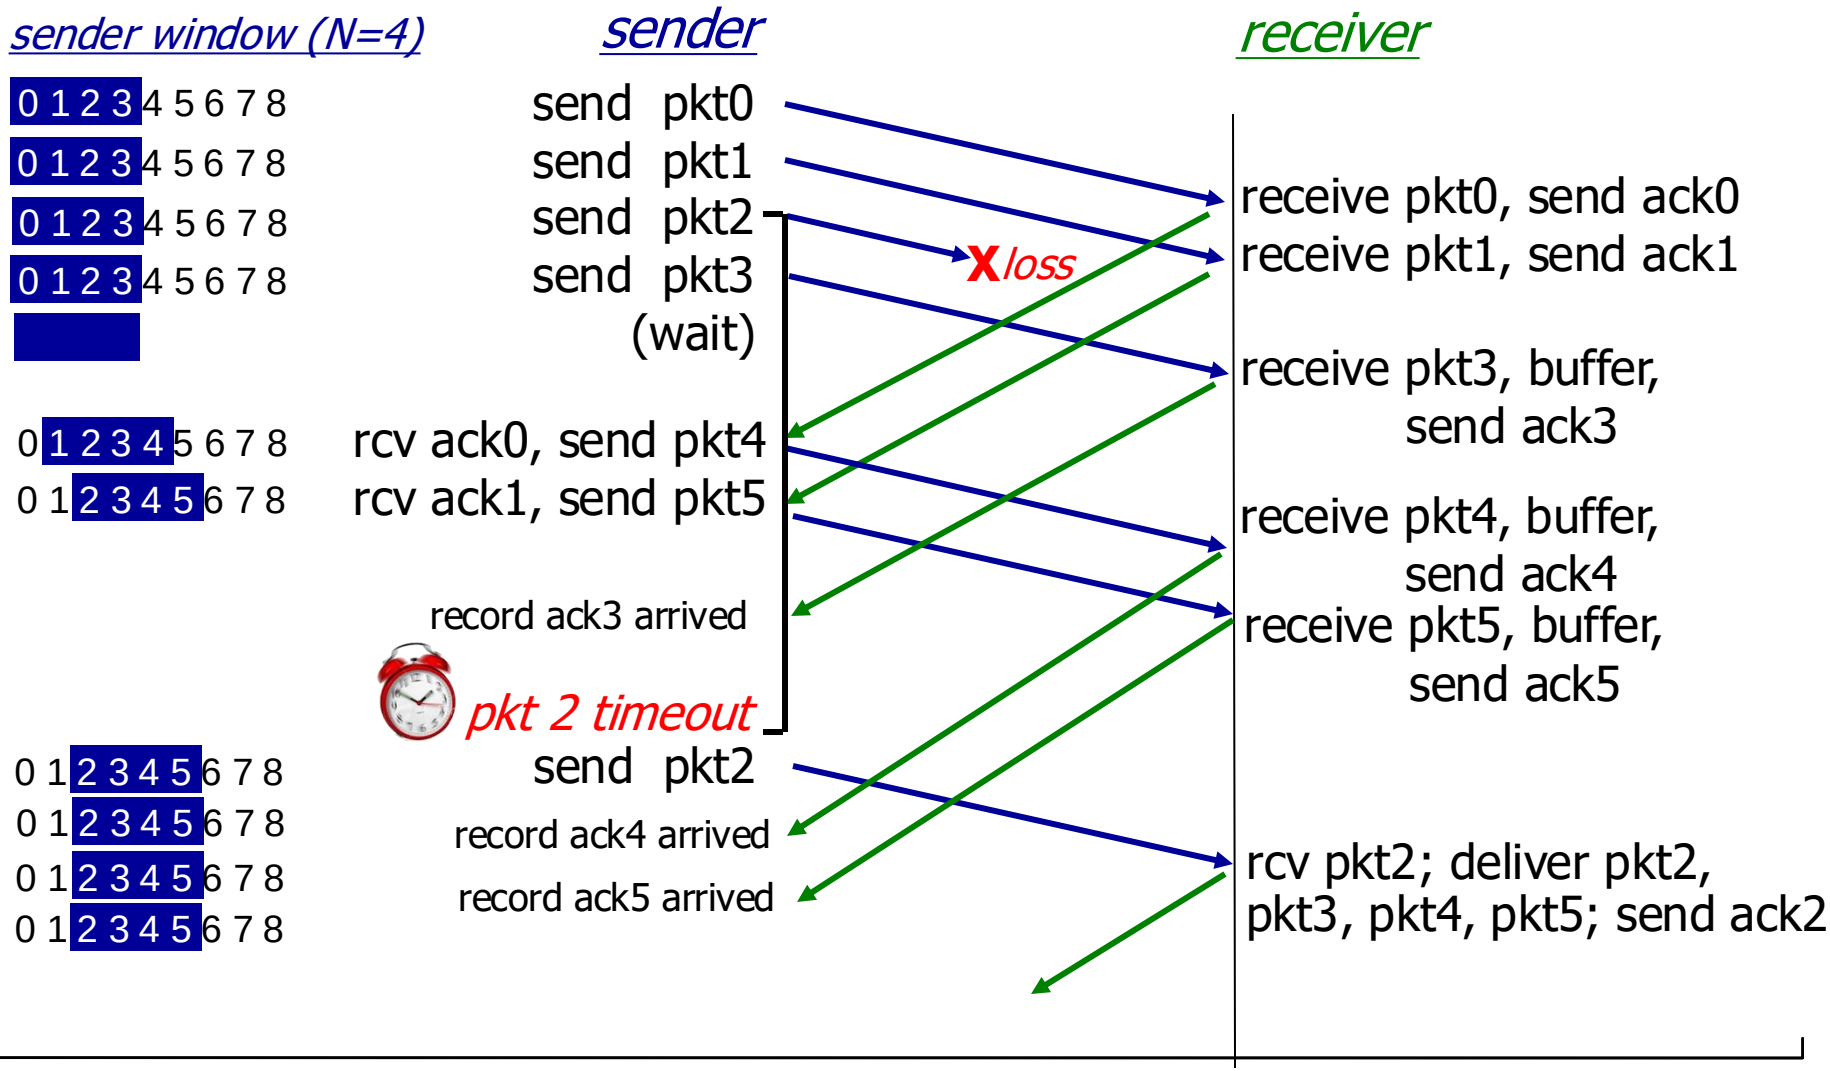

Selective repeat

sender

data from above:

- if next available seq # in window, send pkt

timeout(n):

- resend pkt n, restart timer

ACK(n) in [sendbase,sendbase+N]:

- mark pkt n as received
- if n smallest unACKed pkt, advance window base to next unACKed seq #

receiver

pkt n in [rcvbase, rcvbase+N-1]

- send ACK(n)
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

pkt n in [rcvbase-N,rcvbase-1]

- ACK(n)

otherwise:

- ignore

TCP: Overview RFCs: 793,1122,1323, 2018, 2581

- point-to-point:
  - one sender, one receiver
- reliable, in-order *byte steam:*
  - no "message boundaries"
- pipelined:
  - TCP congestion and flow control set window size
- full duplex data:
  - bi-directional data flow in same connection
- connection-oriented:
  - handshaking (exchange of control msgs) inits sender, receiver state before data exchange
- flow controlled:
  - sender will not overwhelm receiver

TCP segment structure

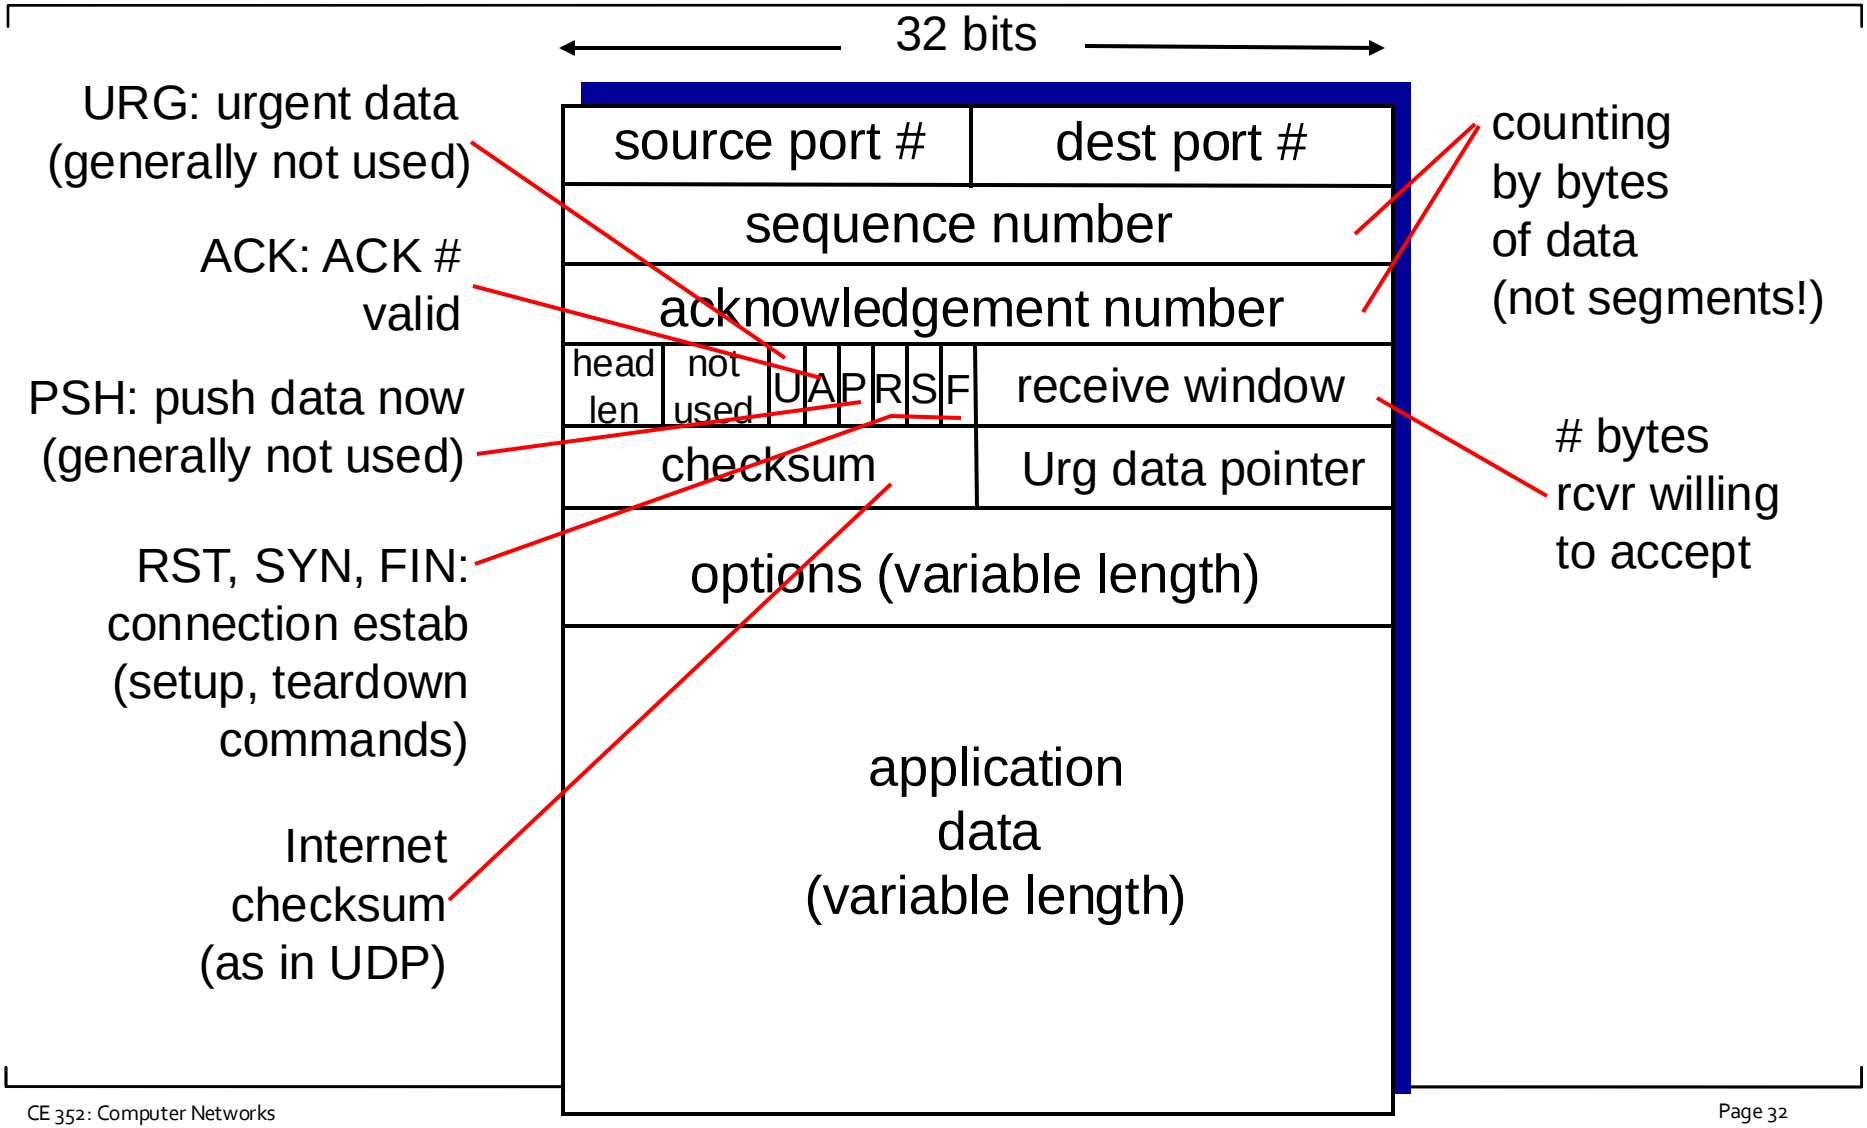

TCP seq. numbers, ACKs

- Sequence numbers:
  - byte stream "number" of first byte in segment's data
- Acknowledgements:
  - seq # of next byte expected from other side
  - cumulative ACK

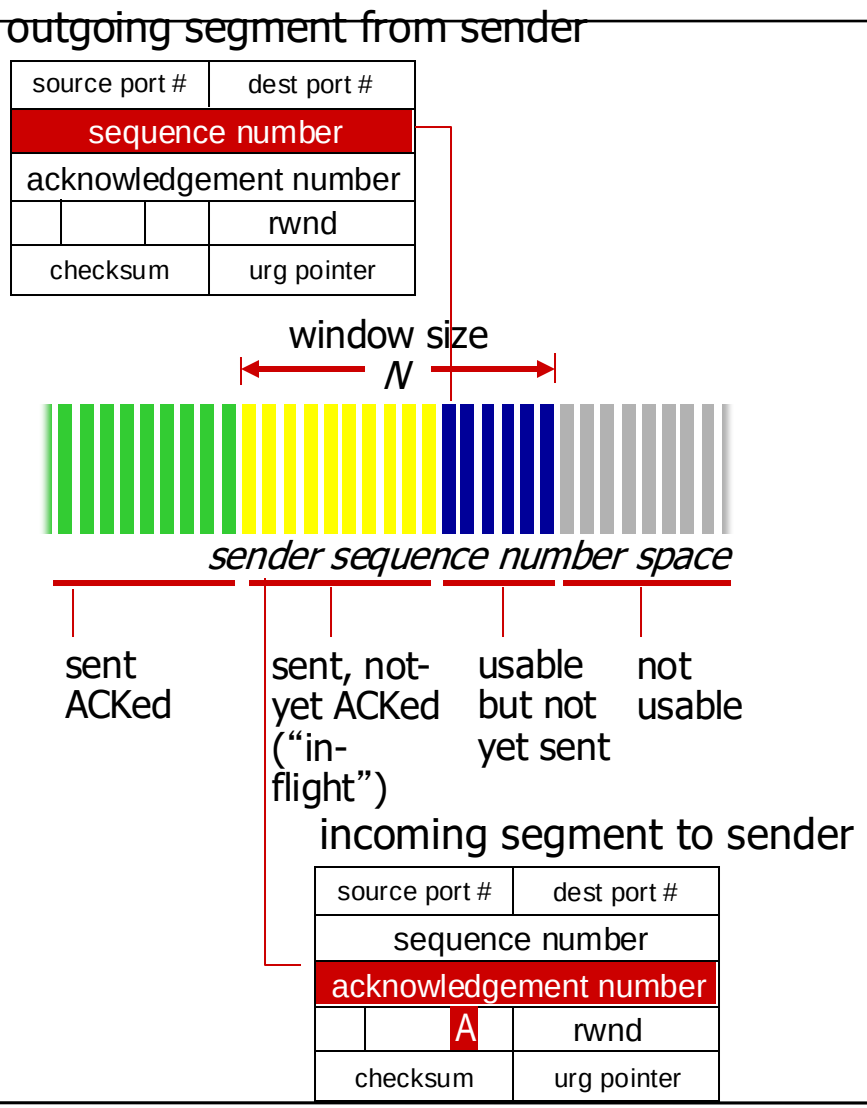

TCP Timeout

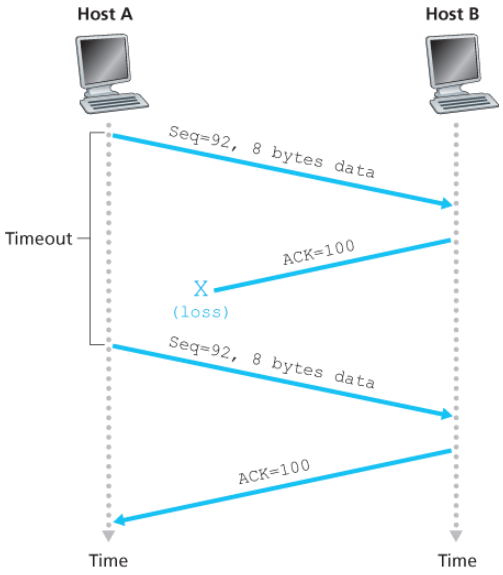

Retransmission due to a lost acknowledgment

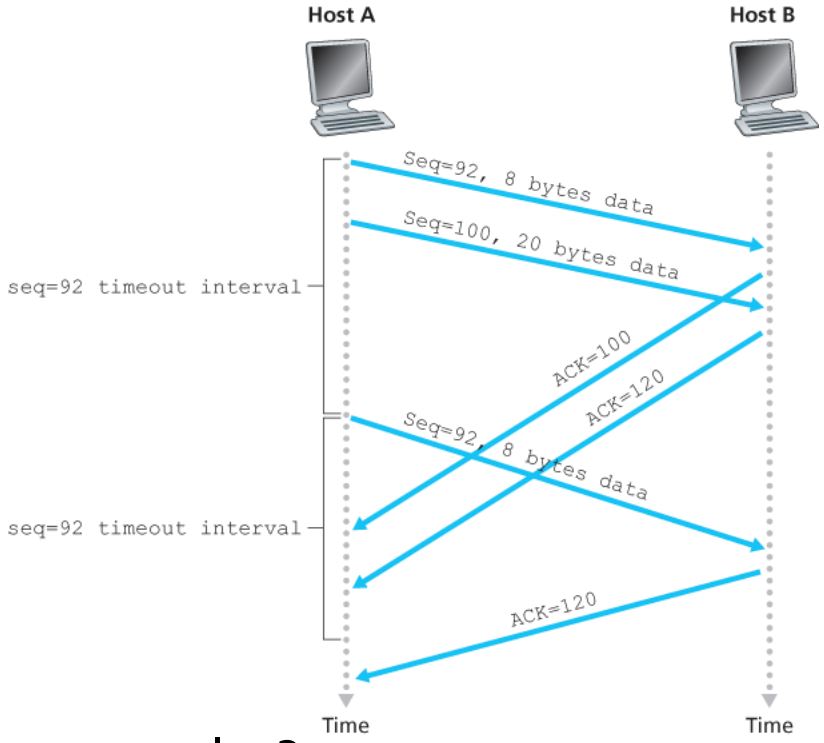

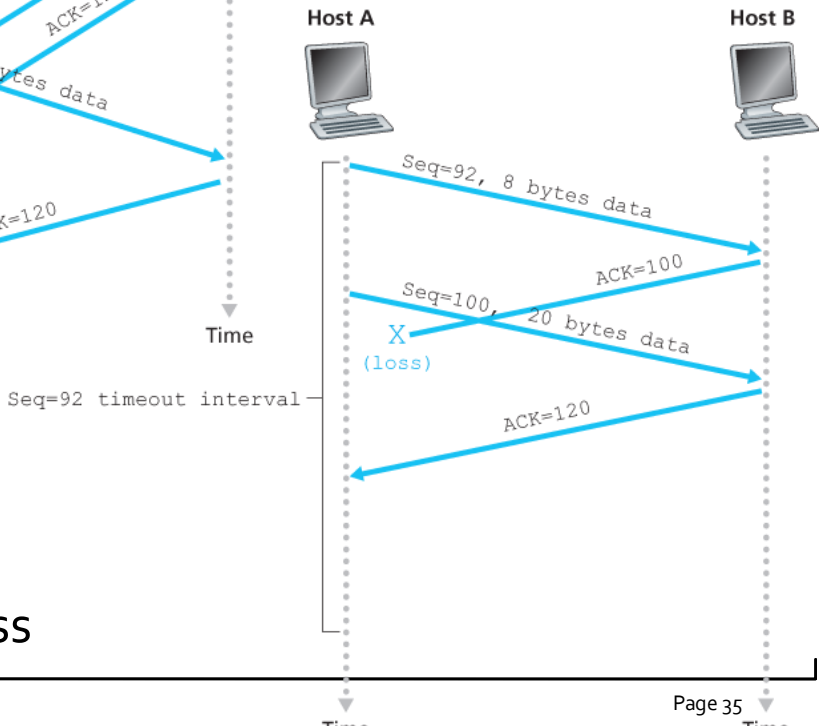

- How to set TCP timeout value?
- longer than RTT, but RTT varies
  - *too short:* premature timeout, unnecessary retransmissions
  - *too long:* slow reaction to segment loss

Bonus 3

- UDP/ IP socket programming
- Client - Server file transfer (connectionless and unreliable)

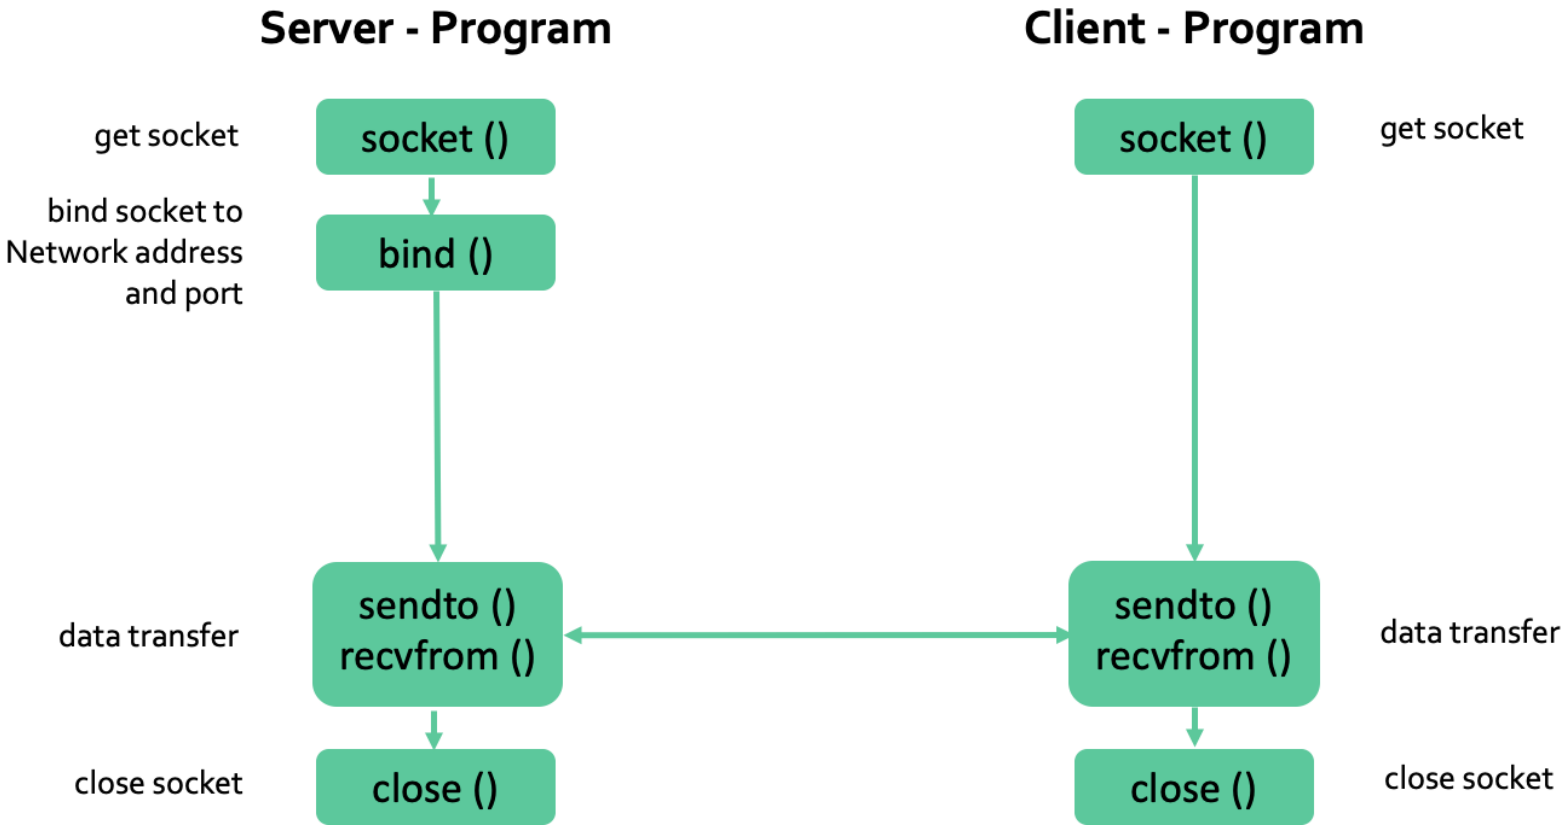

Summary

Today:

- Transport protocols
  - rtd3.0: channel with errors and loss
  - Pipelined protocols: Go-Back-N and Selective repeat
- TCP overview, segment structure, communication

Canvas discussion:

- Reflection
- Exit ticket

Next time:

- read 3.6 and 3.7 of K&R
- follow on Canvas! material and announcements

Any questions?